03 - Core Spark

Demi Ben-Ari

Spark Languages

Spark Word Count example - Spark Shell

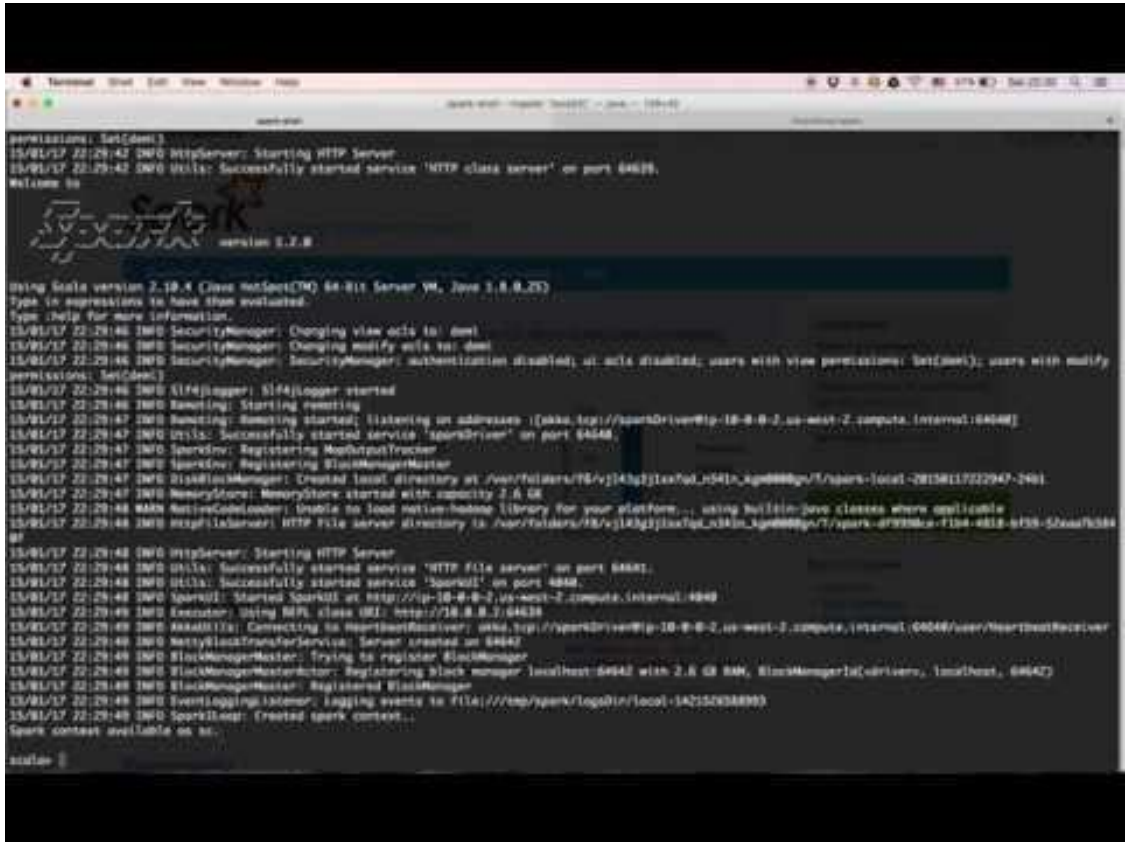

Panorays

Module Overview

- RDD
- Transformations
- Actions

Spark Mechanics

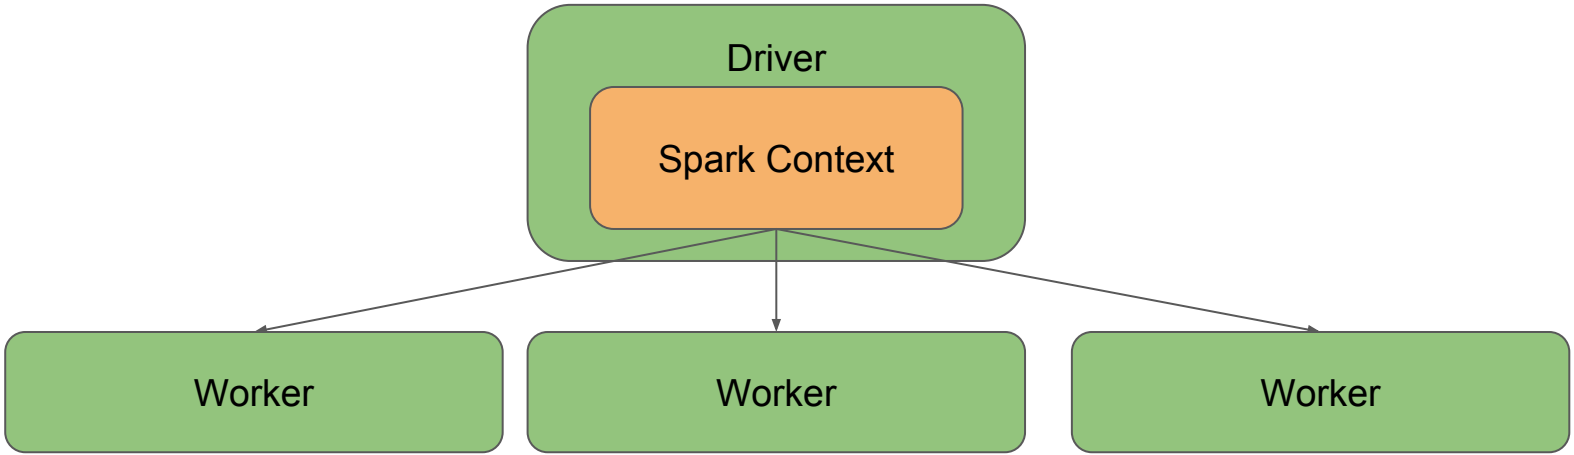

Spark Mechanics

- Task Creator
  - Builds execution graph to be sent to each worker
- Scheduler
  - Scheduling all of the task across the nodes
- Data locality
  - Sending the work to the data to avoid moving data across the cluster
- Fault tolerance
  - Monitoring the tasks for any failures so it can trigger the task on a different node

Spark Application Configuration

Priority and Hierarchy (From most final to the most general):

1) Code
2) spark-submit --Flags
3) app.properties file ([app])
4) defaults (spark-default.sh)

spark-submit example

```
$SPARK_HOME/bin/spark-submit --class org.apache.spark.examples.SparkPi \
    --master yarn-cluster \
    --num-executors 10 \
    --executor-cores 2 \
    spark-examples-1.6.0-hadoop2.6.0.jar \
    100
```

RDD - Resilient Distributed Dataset

- … Collection of elements partitioned across the nodes of the cluster that can be operated on it in parallel…
  - http://spark.apache.org/docs/latest/programming-guide.html#overview
- RDD - Resilient Distributed Dataset
  - Collection similar to a List / Array (Abstraction)
  - It’s actually an Interface (Behind the scenes it’s distributed over the cluster)
- DAG - Directed Acyclic Graph
- Are Immutable!!!

RDD - Resilient Distributed Dataset

- Transformations are Lazy evaluated
  - map
  - filter
  - .....
- Actions - Triggers DAG computation
  - collect
  - count
  - reduce

What's really an RDD???

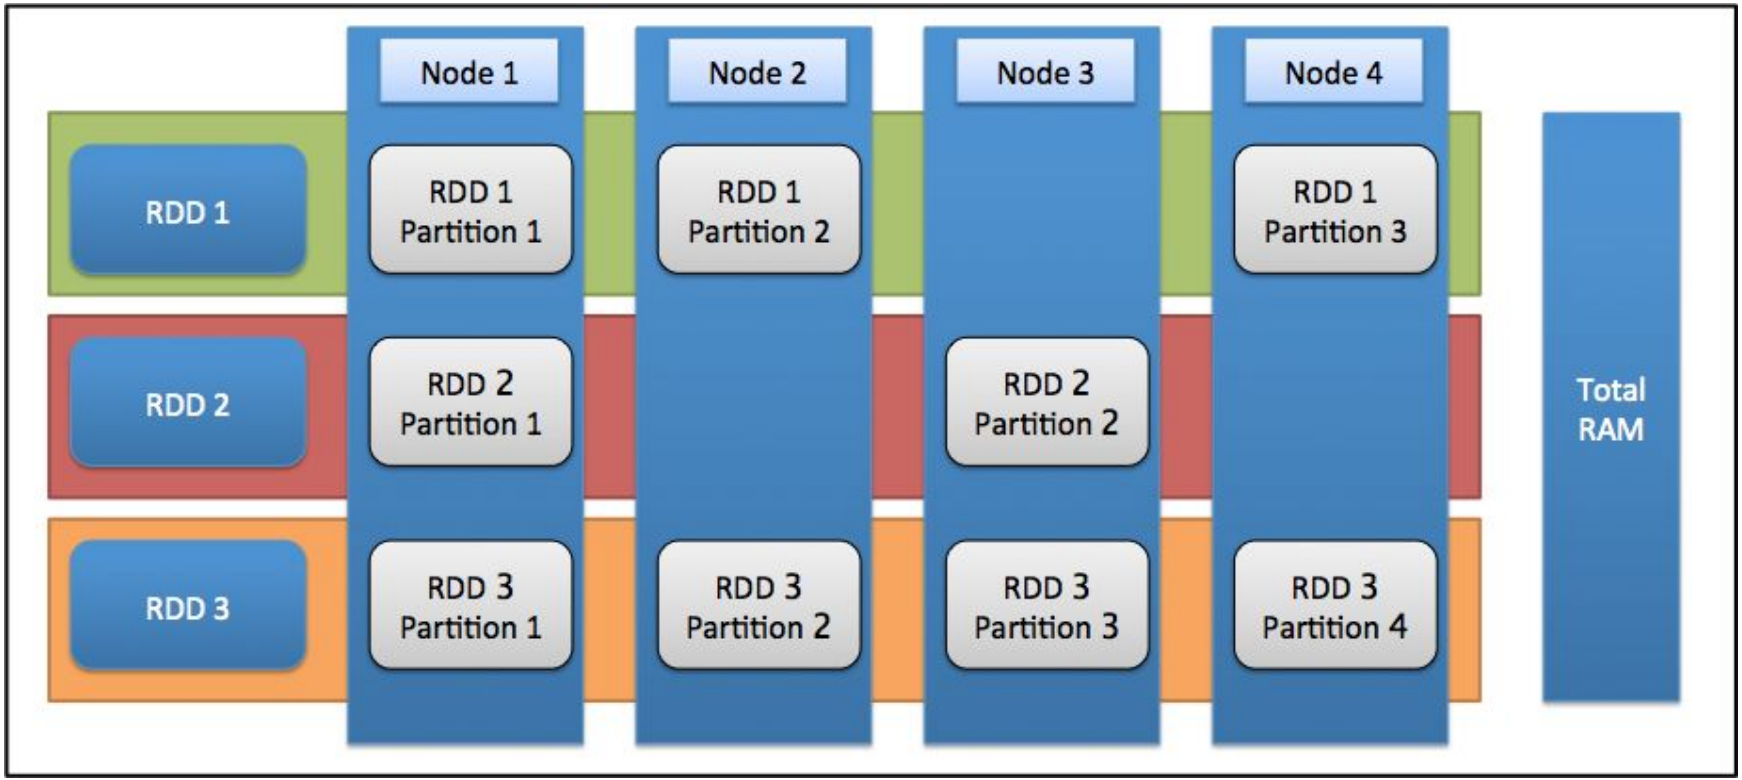

Panorays

Spark Mechanics

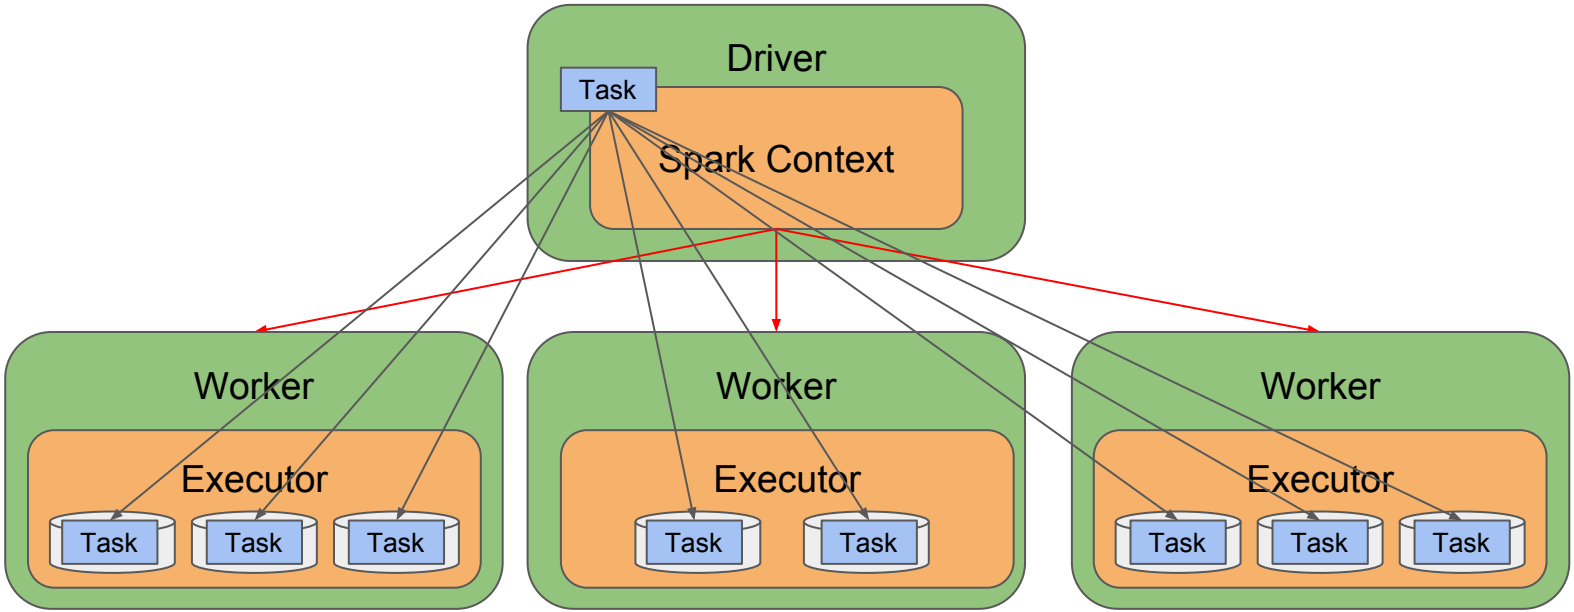

Panorays

Lambdas - Anonymous Functions

- How It’s done in Java
- Why not to use it :)
  - Not testable
  - Results in a very verbose main class (inner classes)

Input

Input methods

- Local FileSystem
- HDFS
- Cassandra
- Avro
- Parquet

Input methods

- sc.parallelize(<Collection>, <number of partitions>)
- sc.textFile(<path>)
- sc.sequenceFile()
  - Hadoop format
- sc.newAPIHadoopFile()
  - instead of partitioning it accepts Hadoop configuration

Transformations

Thanks to Jeffrey Thompson for many of the schemas
https://github.com/jkthompson/pyspark-pictures

Official Documentation

Wide and Narrow Transformations

- **Narrow dependency**: each partition of the parent RDD is used by at most one partition of the child RDD. This means the task can be executed locally and we **don't have to shuffle.** (Eg: map, flatMap, Filter, sample)
- **Wide dependency**: multiple child partitions may depend on one partition of the parent RDD. **This means we have to shuffle data** unless the parents are hash-partitioned (Eg: sortByKey, reduceByKey, groupByKey, cogroupByKey, join, cartesian)
- You can read a good blog post about it.

Basic Terms - Wide and Narrow Transformations

Narrow Dependencies:

filter,
map...

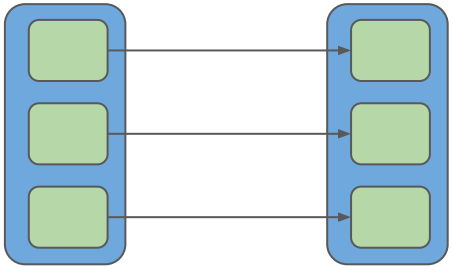

union

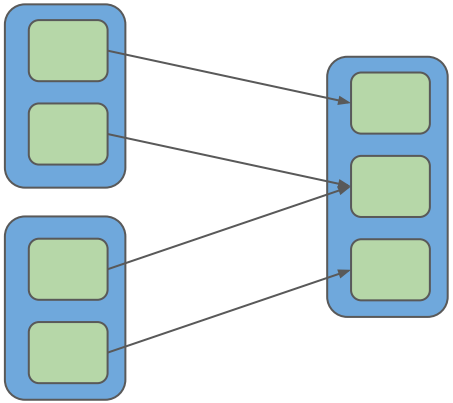

Wide (Shuffle) Dependencies:

groupByKey

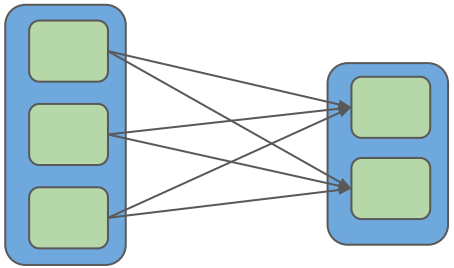

join...

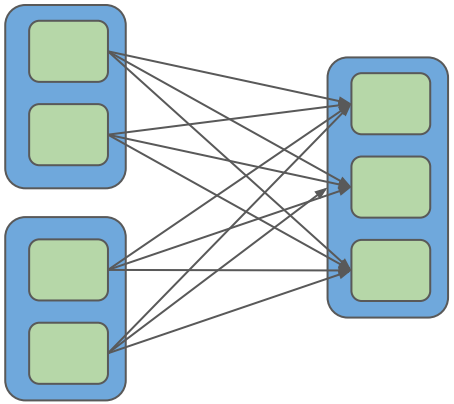

Panorays

RDD

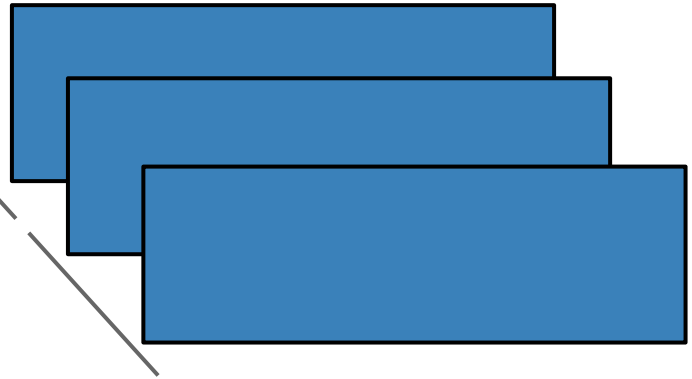

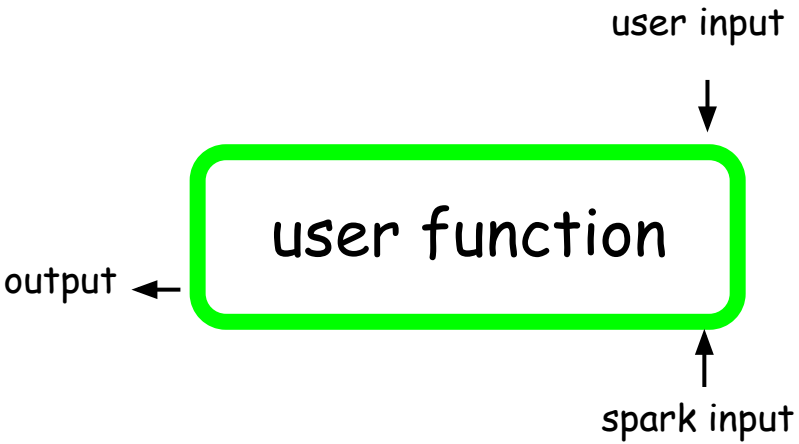

RDD Elements

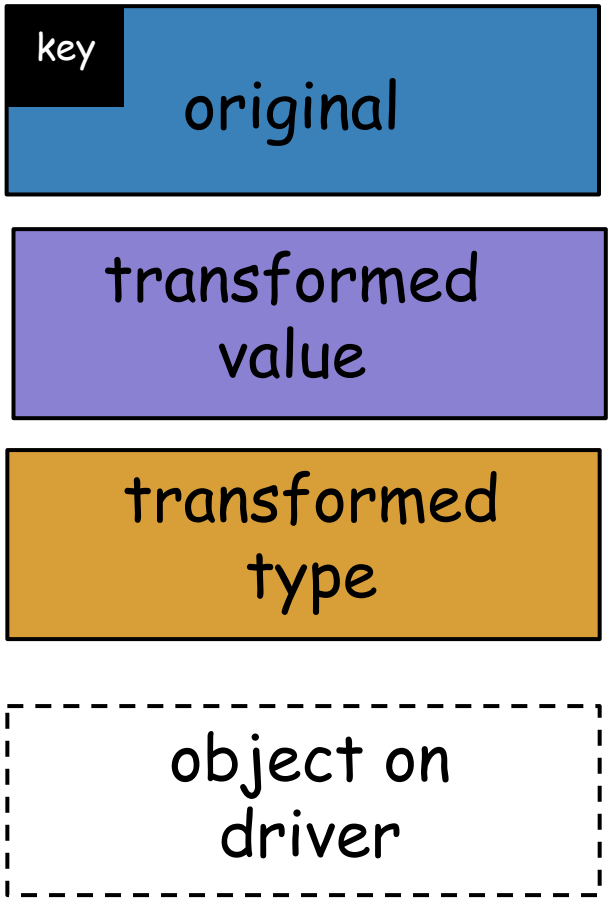

filter

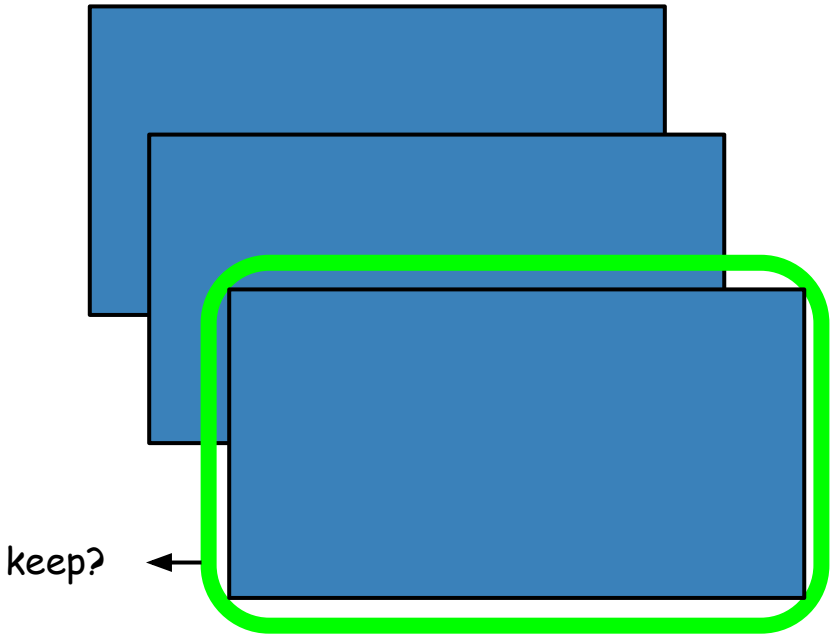

Panorays

map

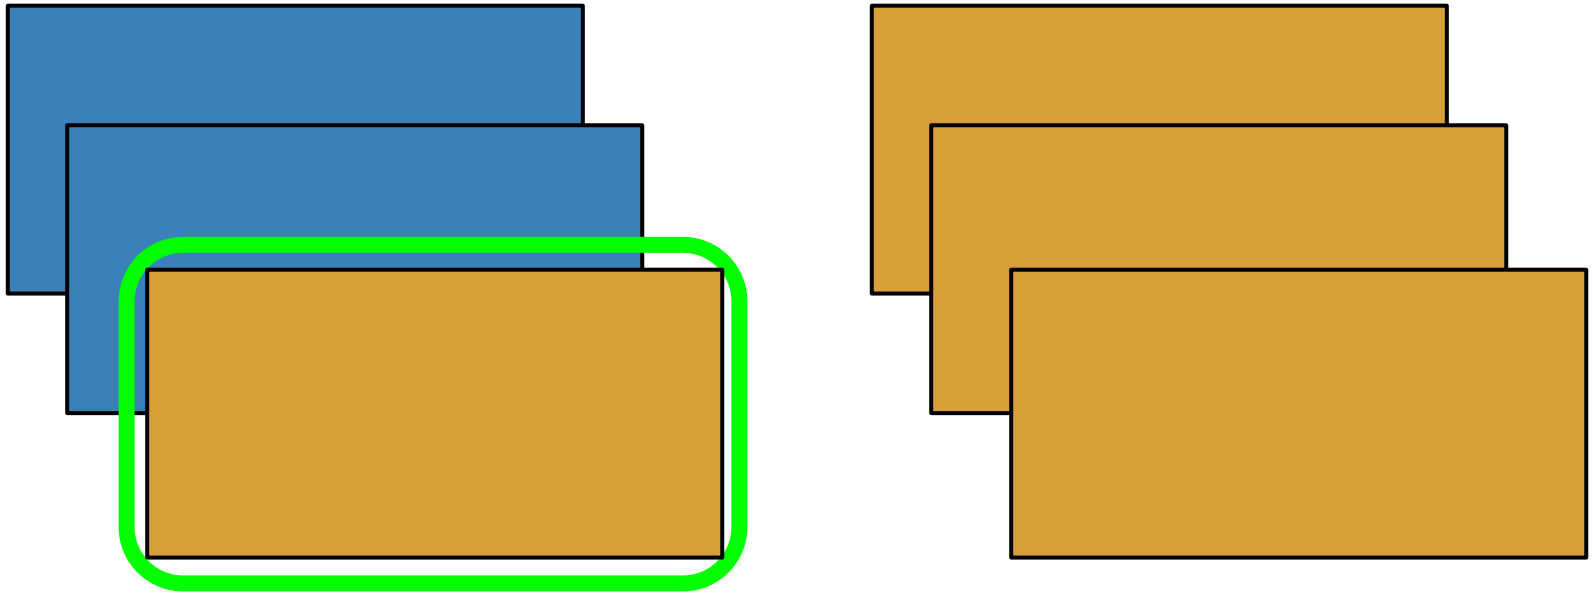

mapValues

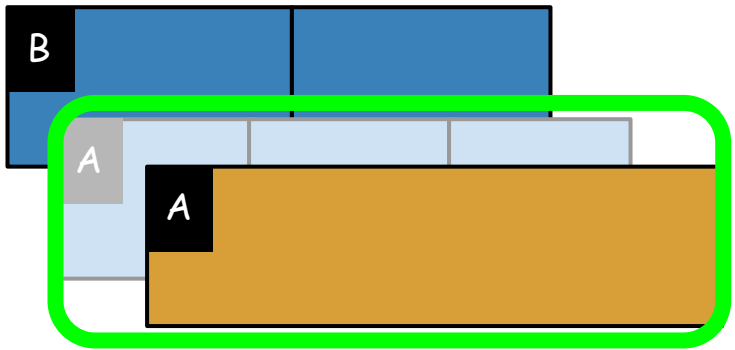

flatMap

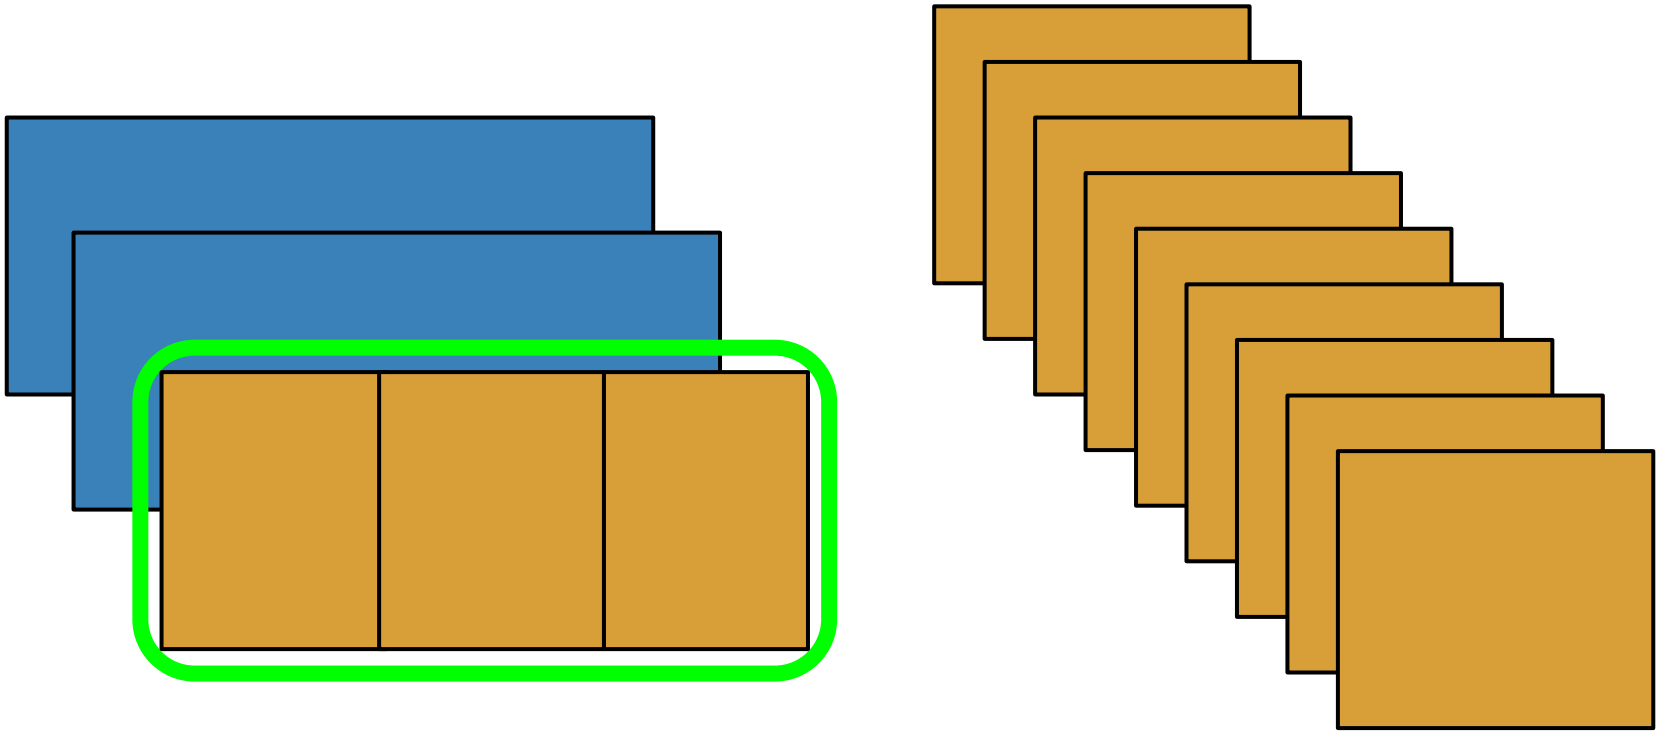

flatMapValue

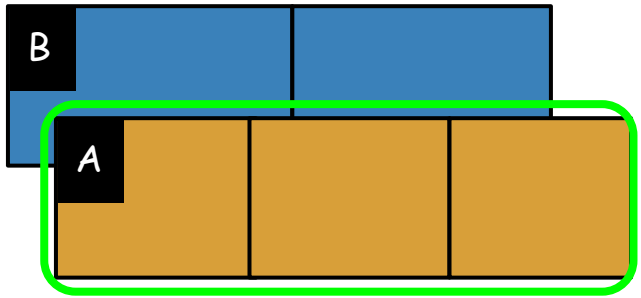

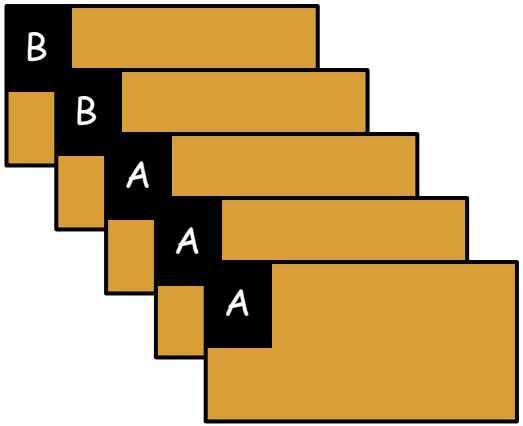

mapPartitions

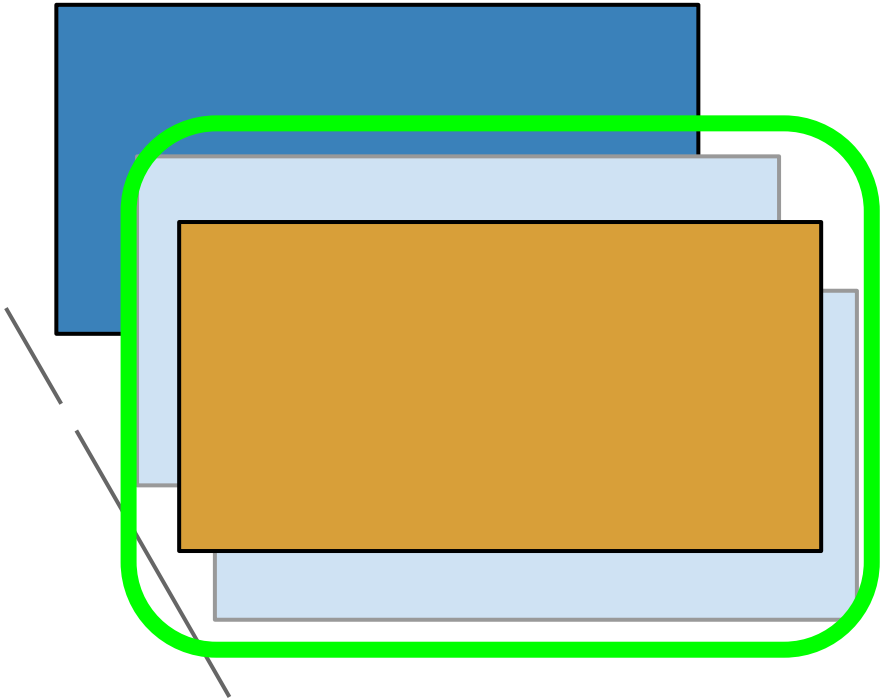

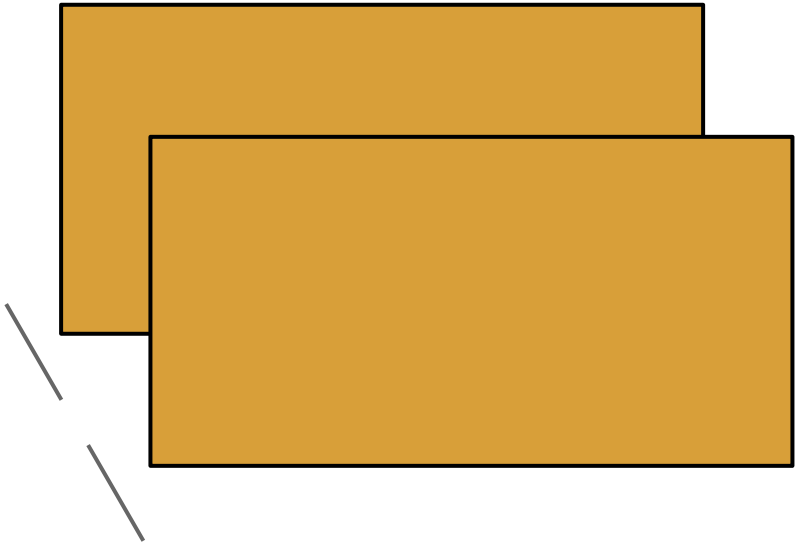

Panorays

mapPartitionsWithIndex

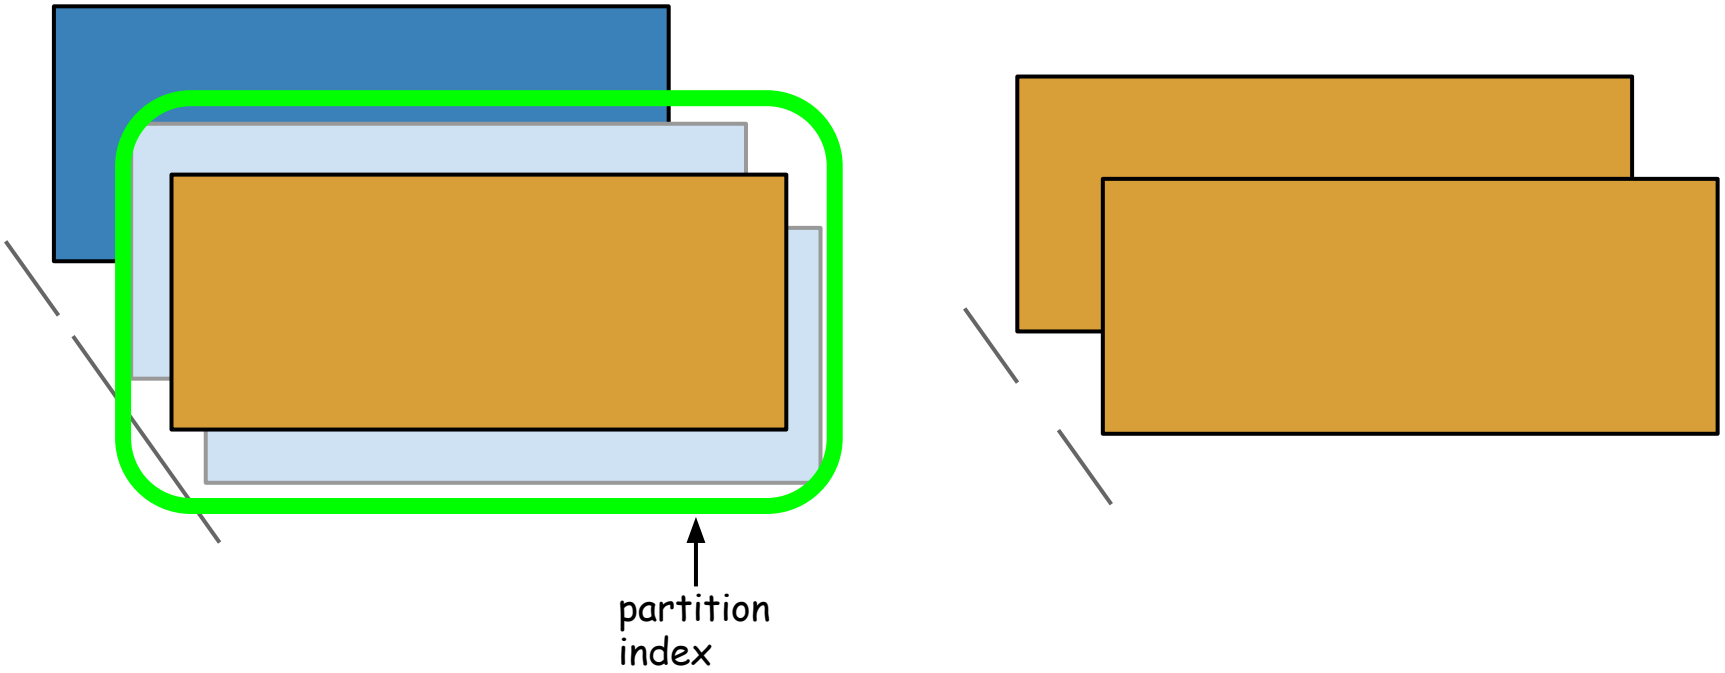

getNumPartitions

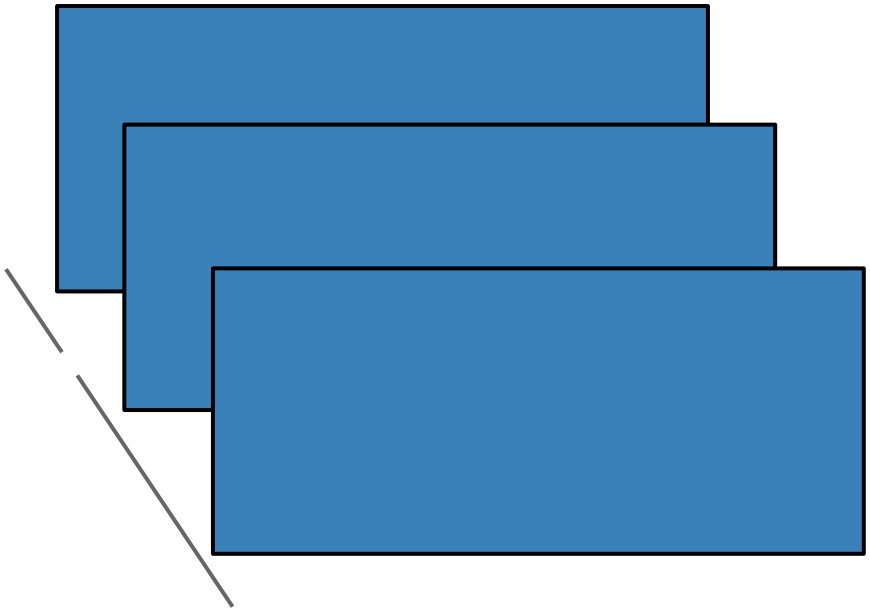

2

distinct

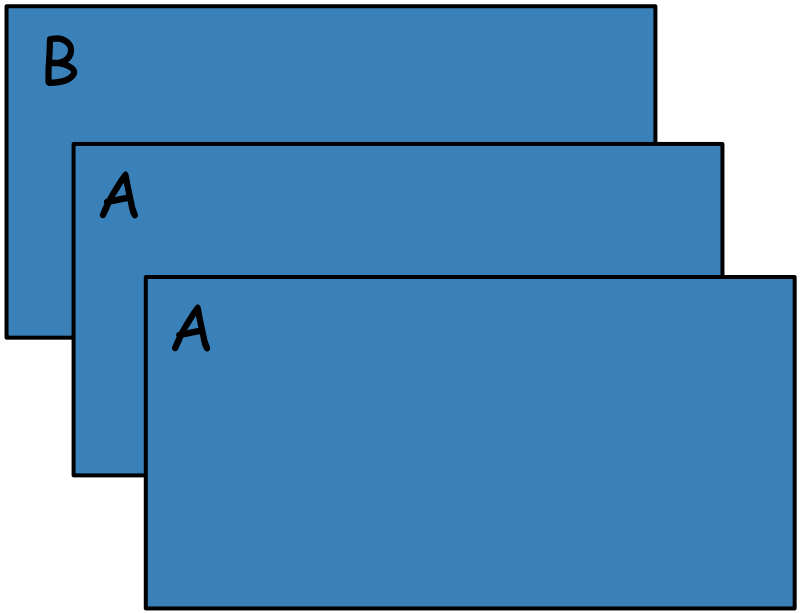

Panorays

sample

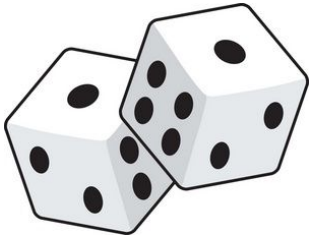

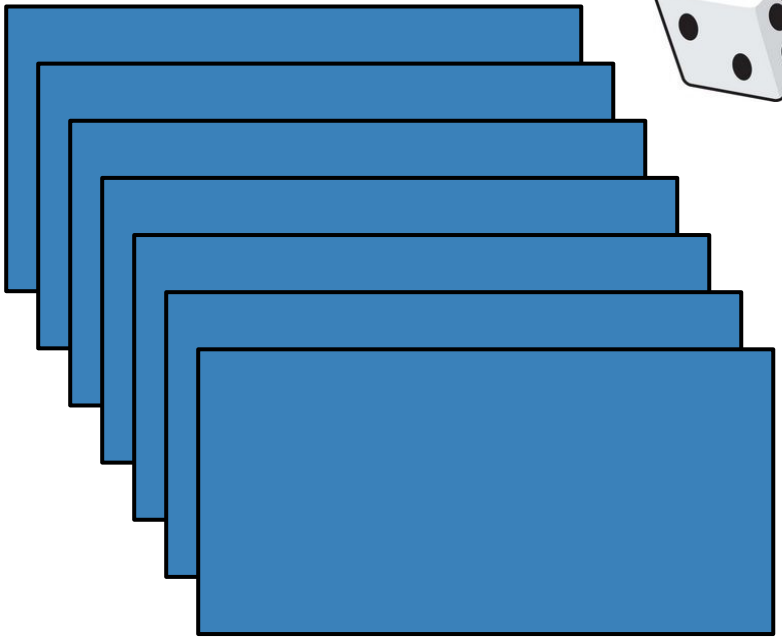

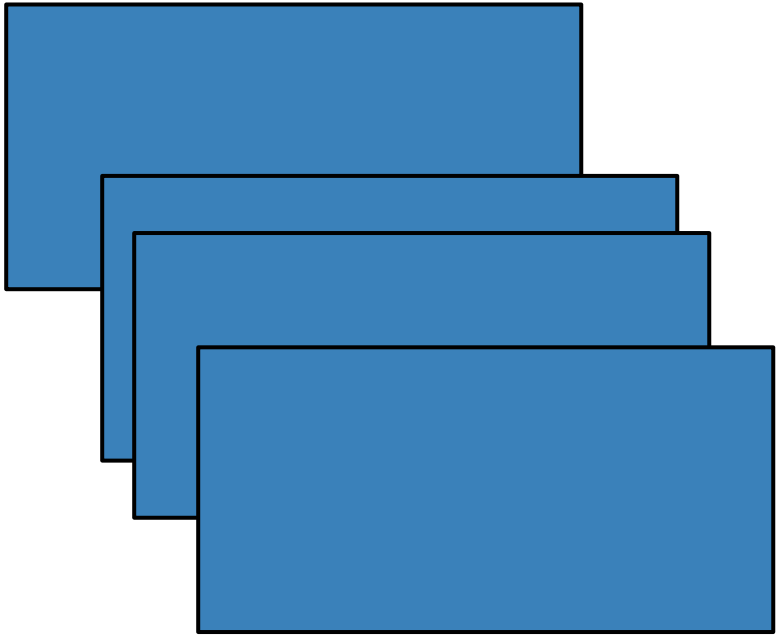

Panorays

glom

- I have never used it :)

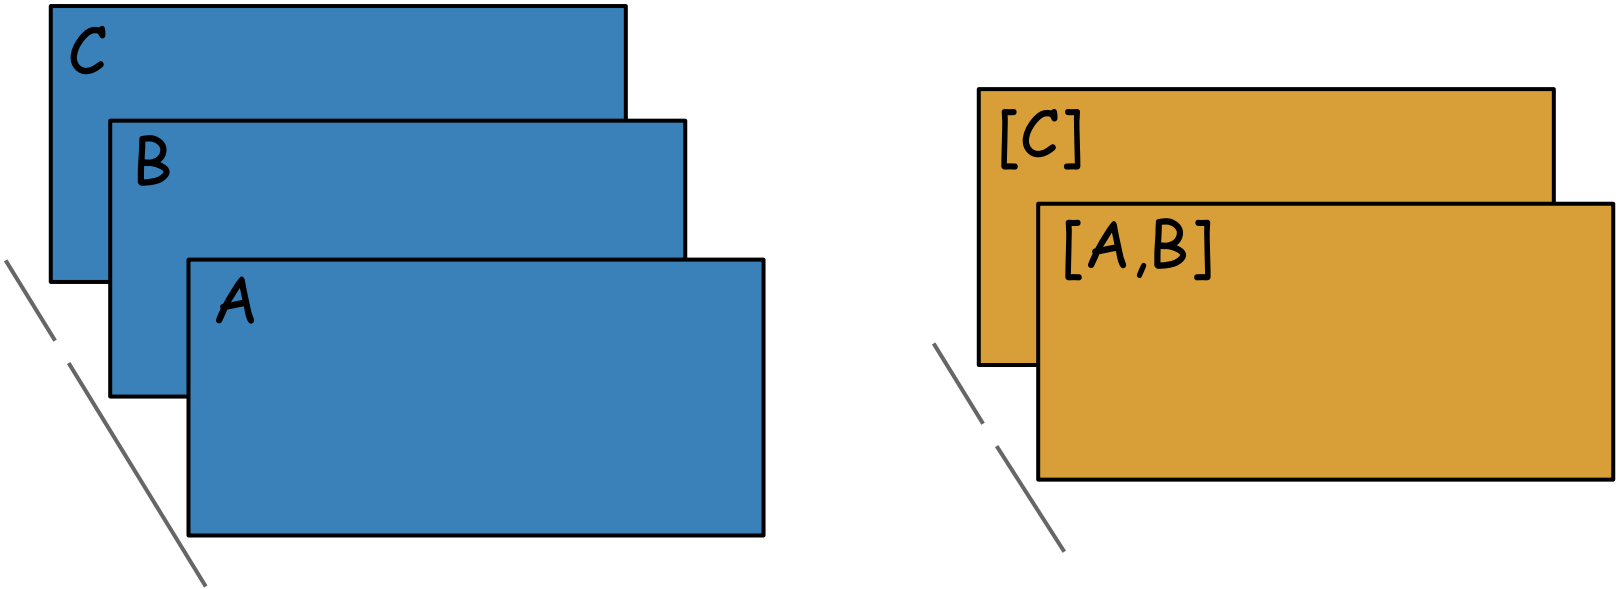

Union (RDD1.union(RDD2))

- Can also use: sc.unionRDDs(RDD1, RDD2, RDD3...)

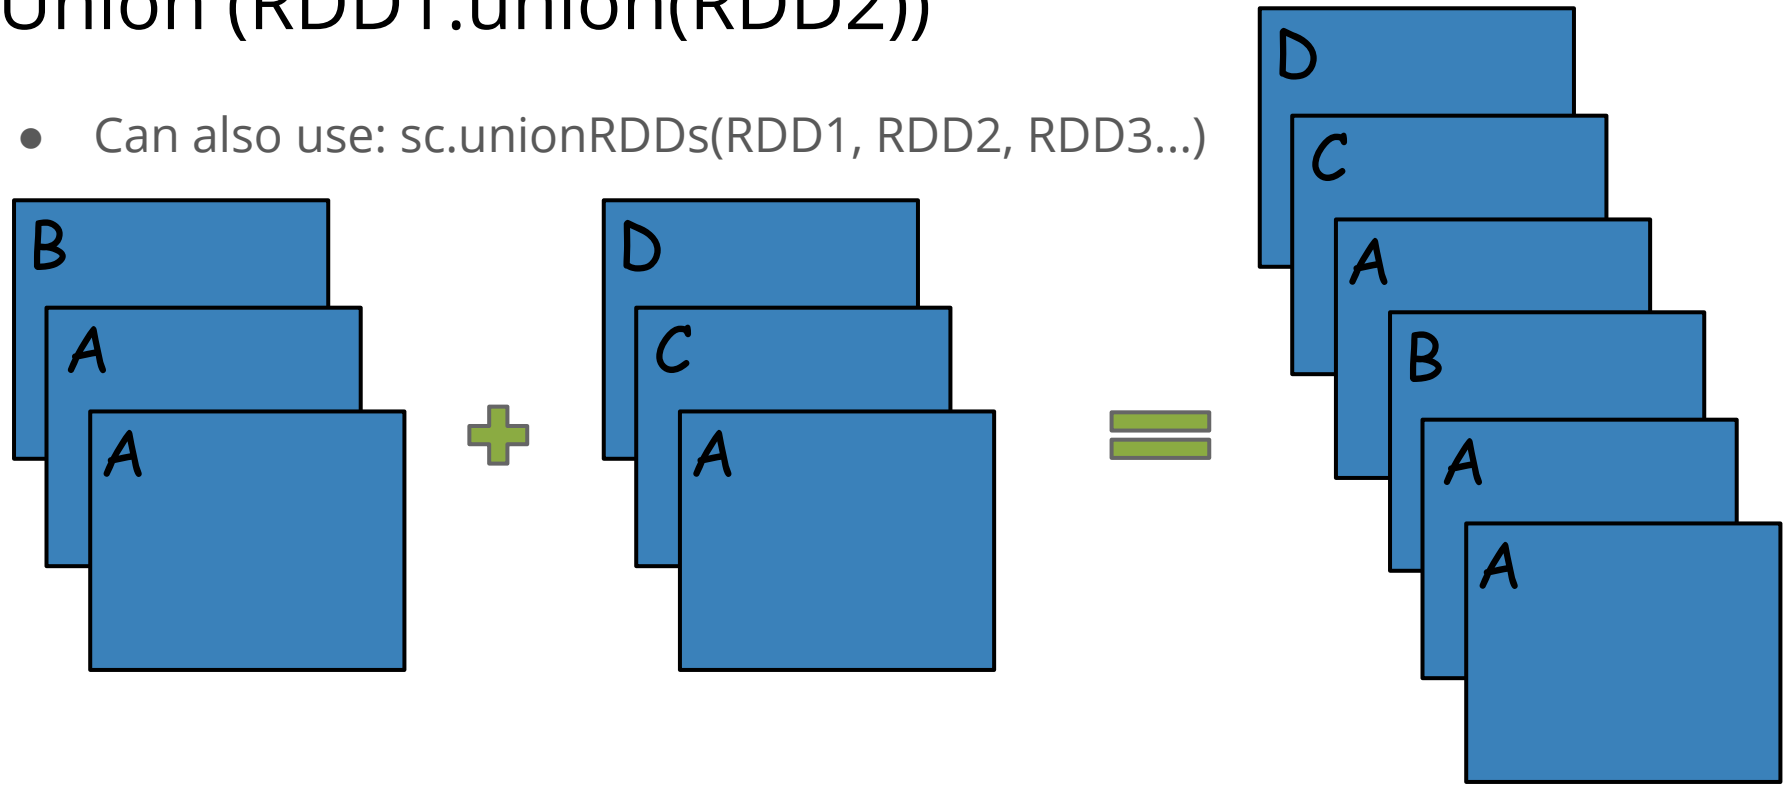

Intersection (RDD1.intersection(RDD2))

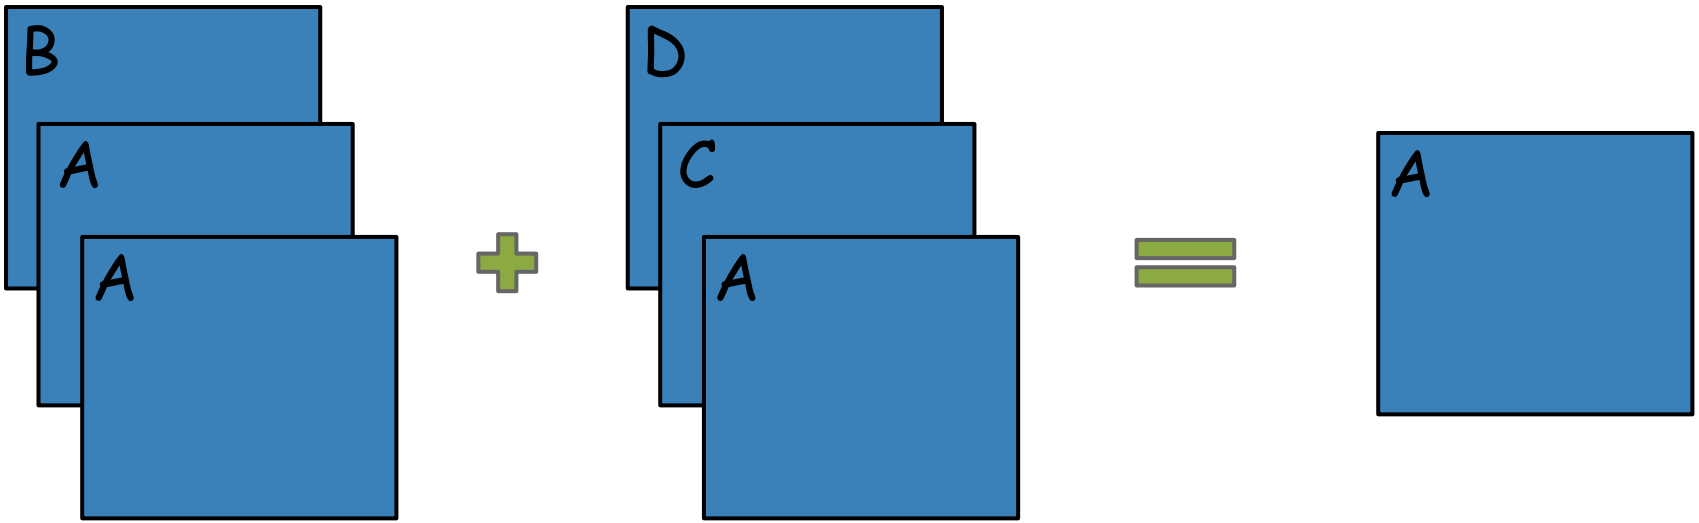

repartition

numPartitions = 3

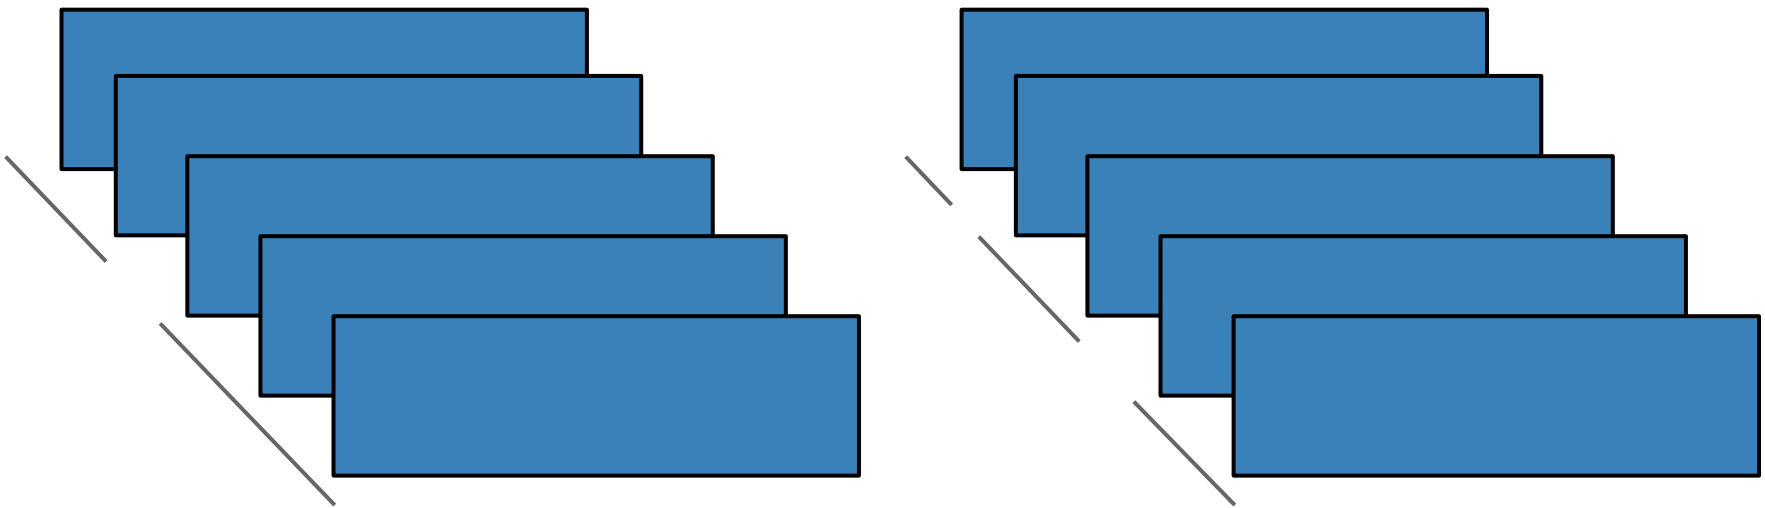

Panorays

coalesce

numPartitions = 1

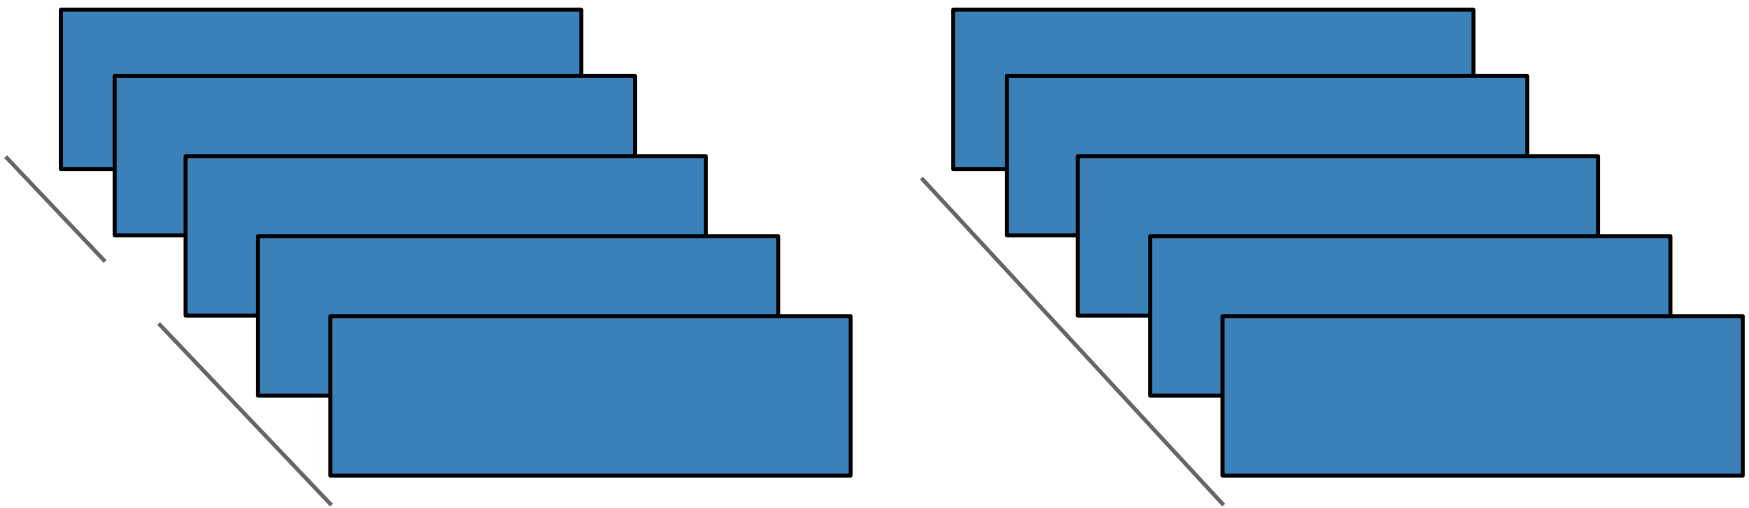

cartesian

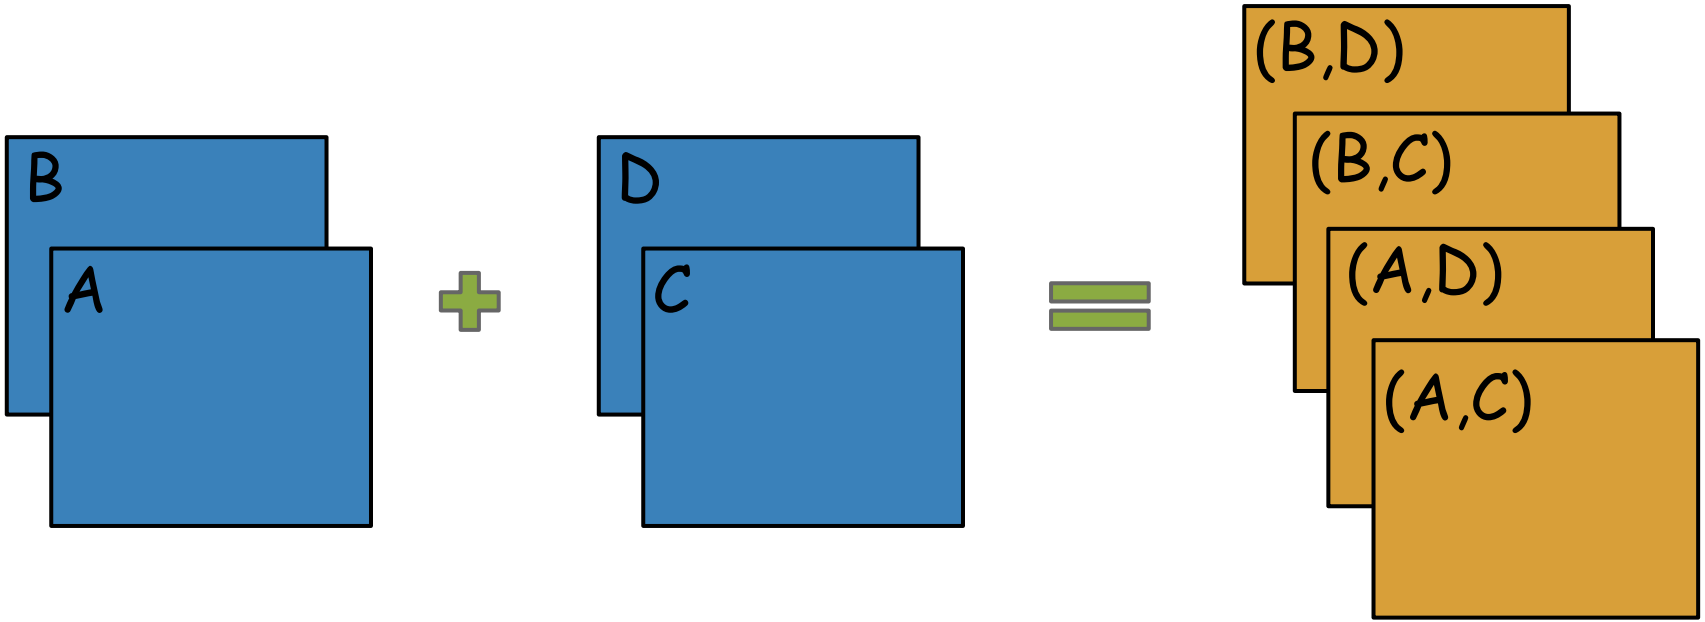

substruct (RDD1.substruct(RDD2))

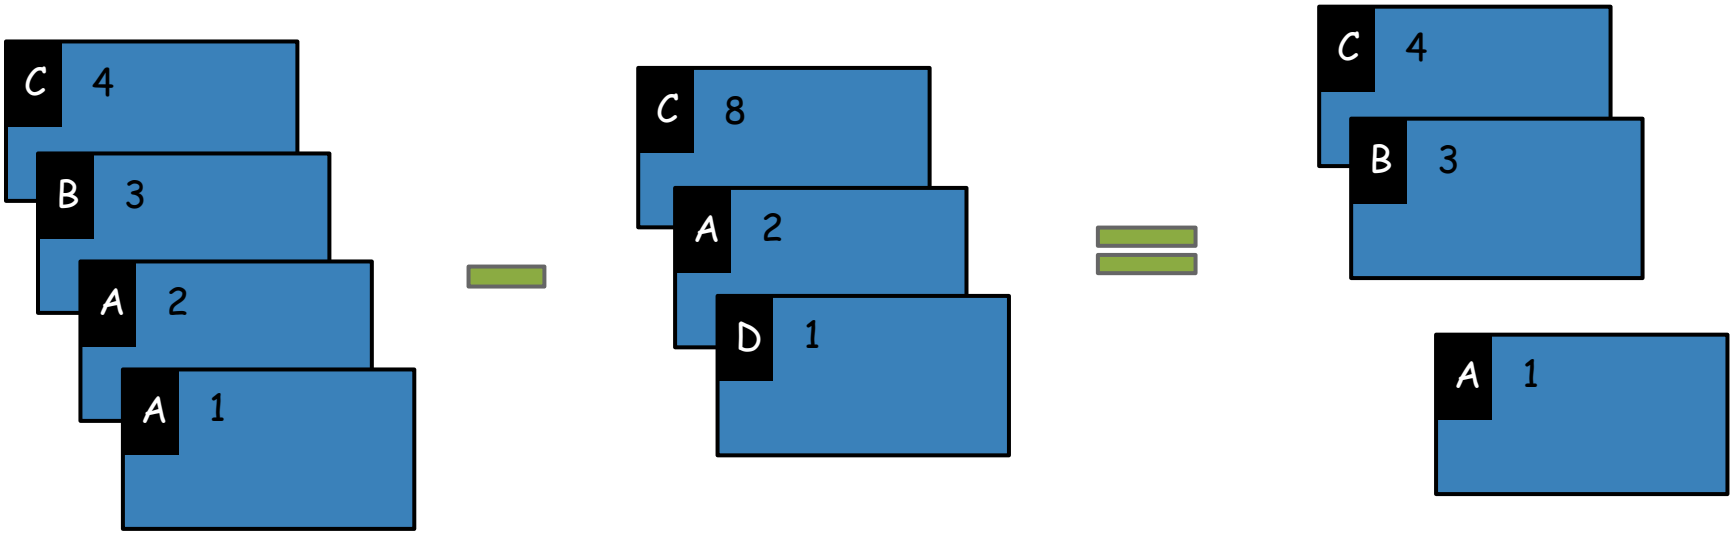

Panorays

groupBy

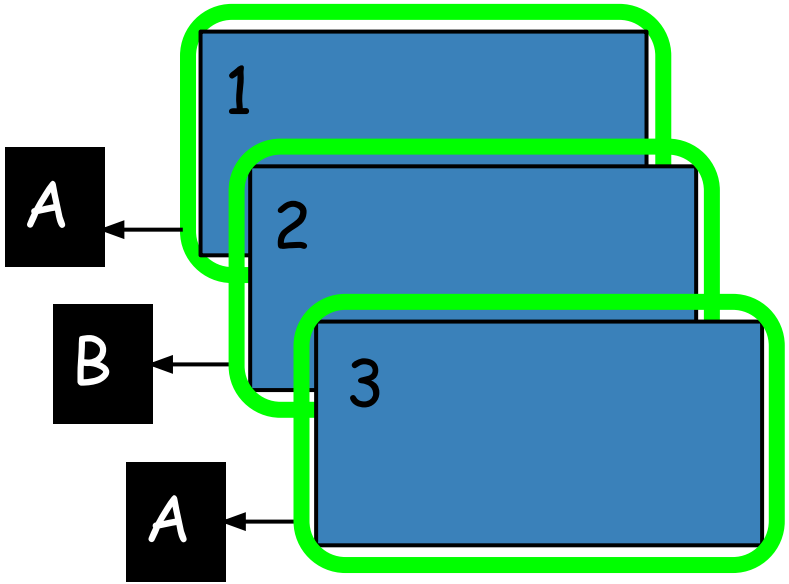

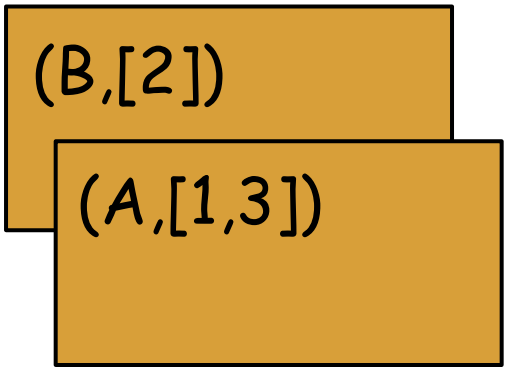

groupByKey

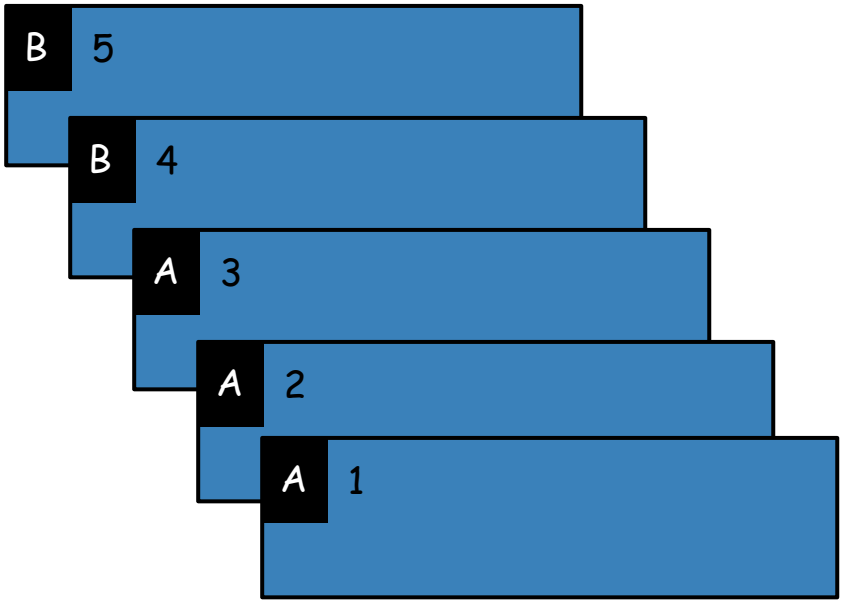

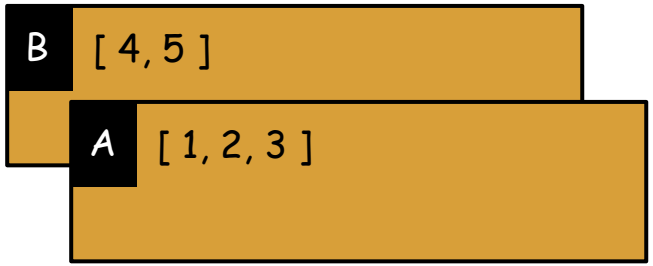

Panorays

cartesian (RDD1.cartesian(RDD2))

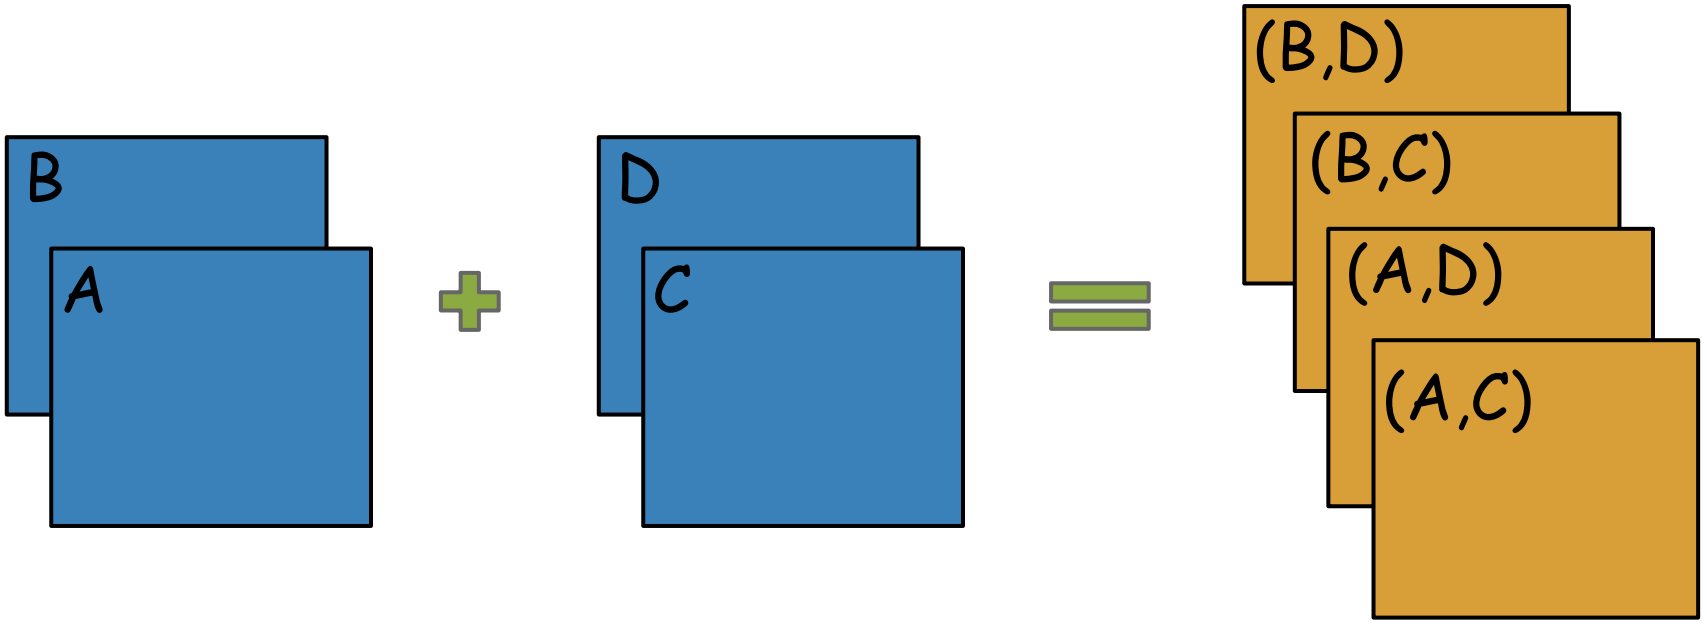

sortByKey

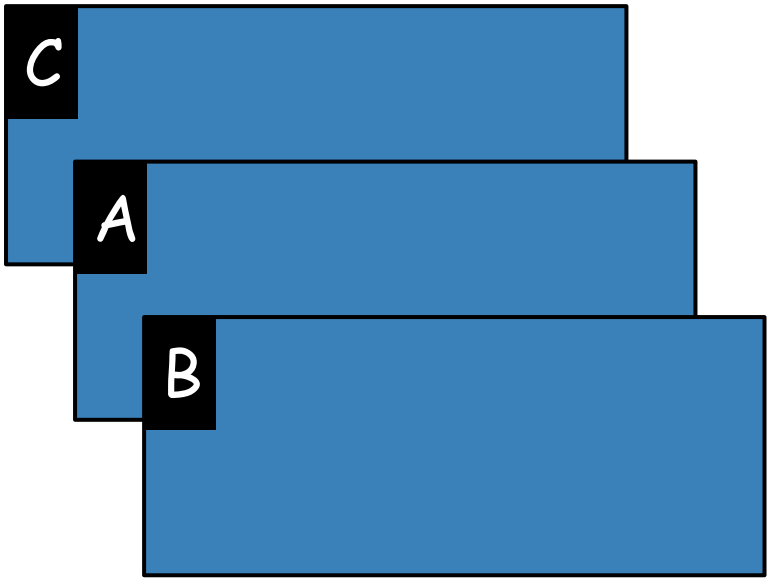

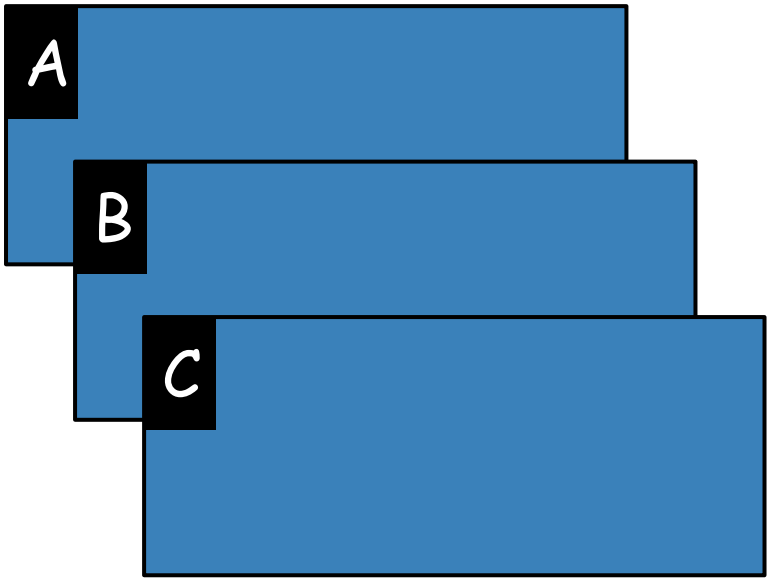

Panorays

sortBy

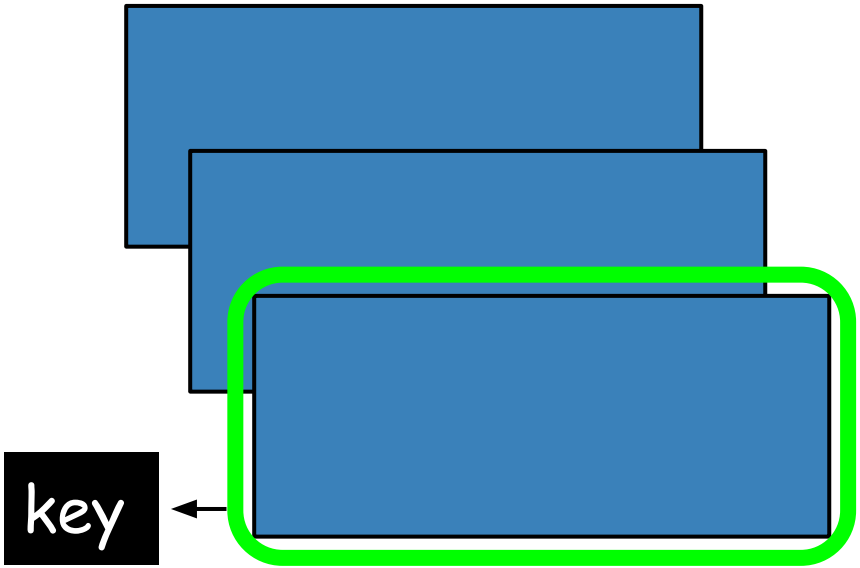

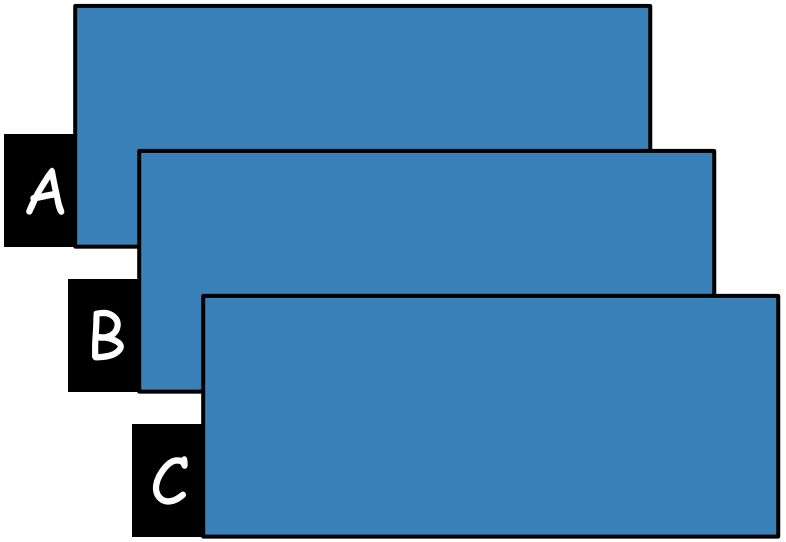

Panorays

pipe

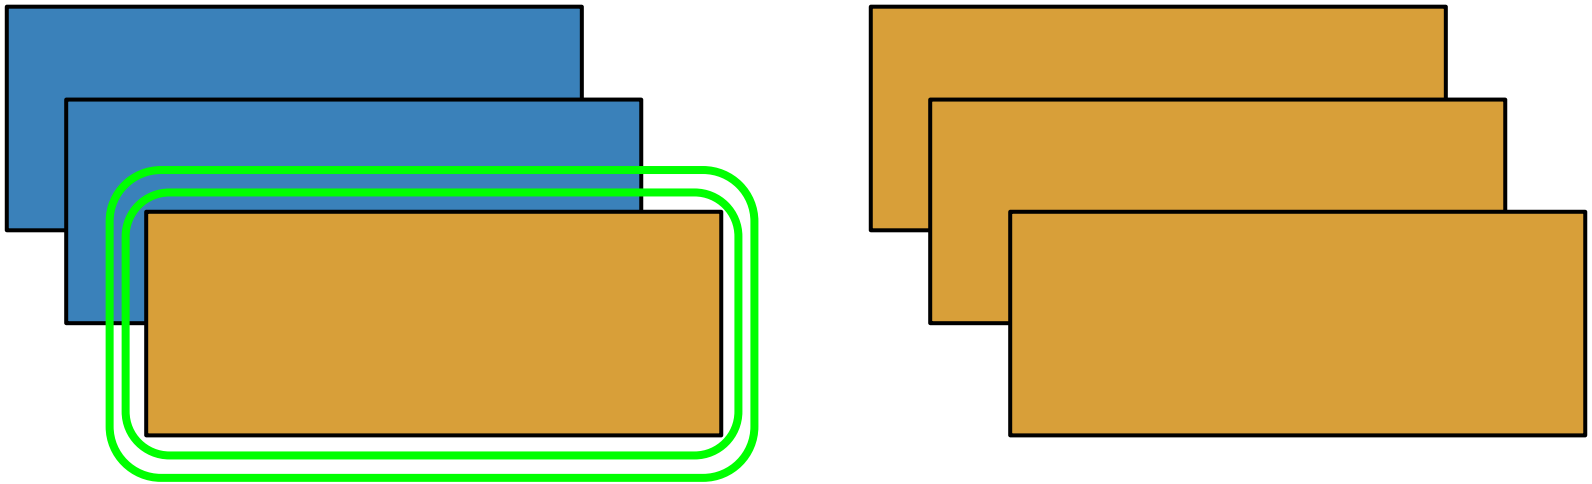

- Pipe each partition of the RDD through a shell command, e.g. a Perl or bash script. RDD elements are written to the process's stdin and lines output to its stdout are returned as an RDD of strings.
- (I've never used it)

zip (RDD1.zip(RDD2))

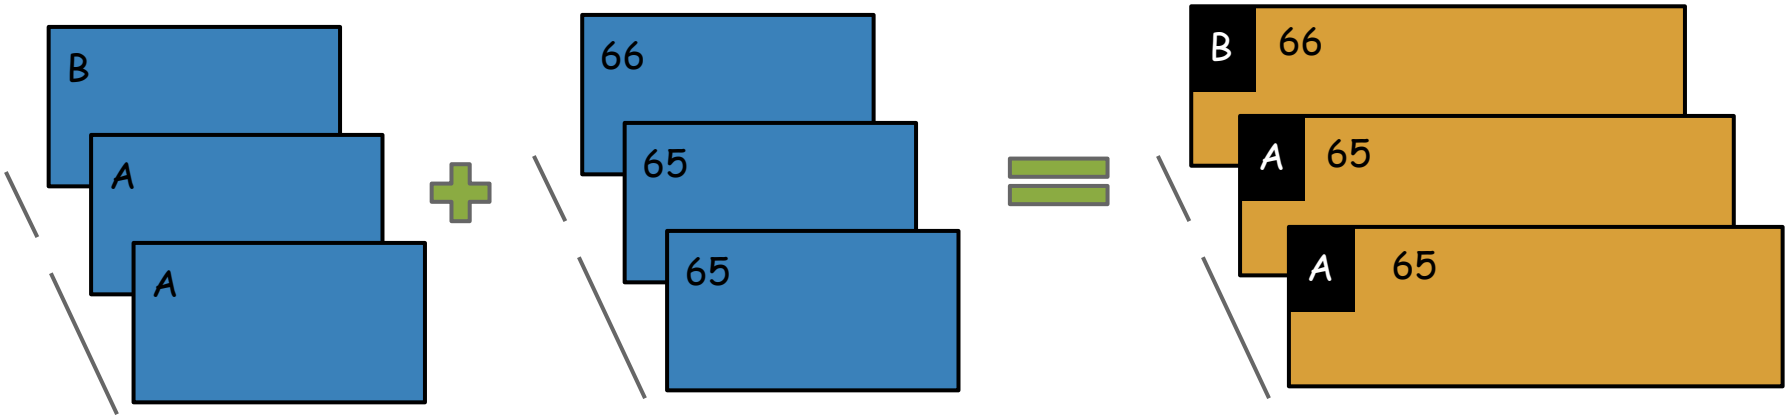

Actions

count

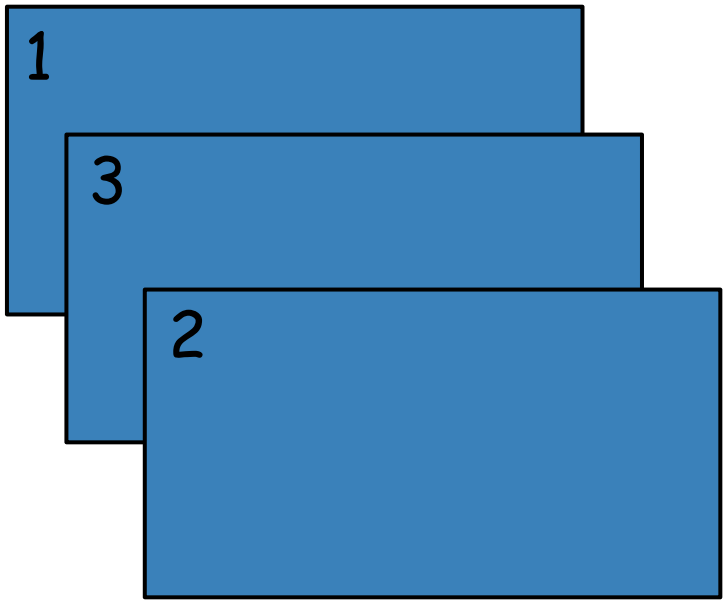

takeSample

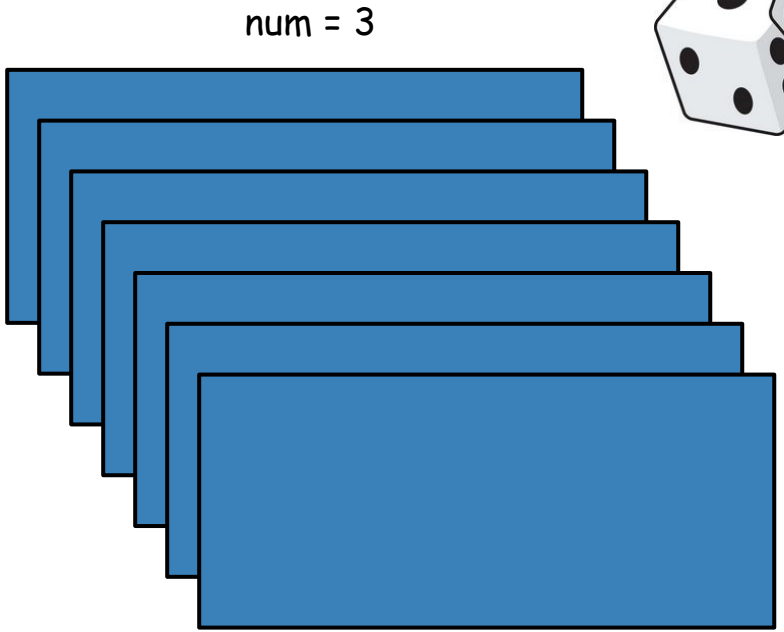

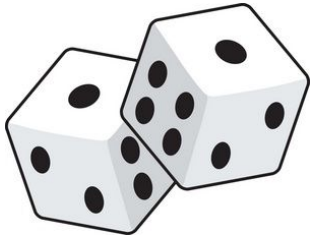

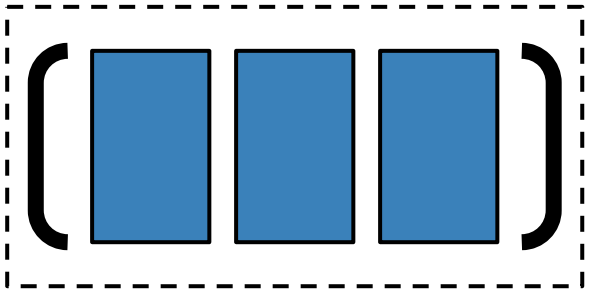

Collect

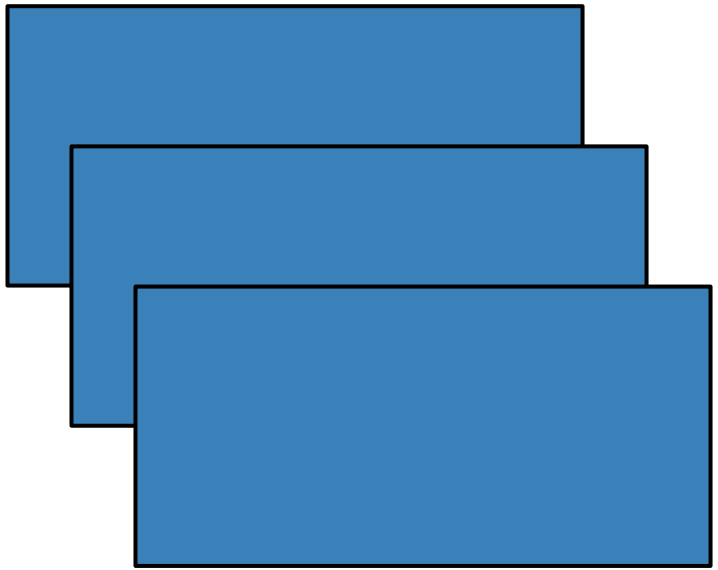

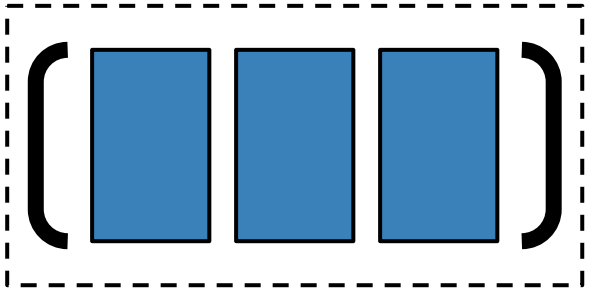

foreach

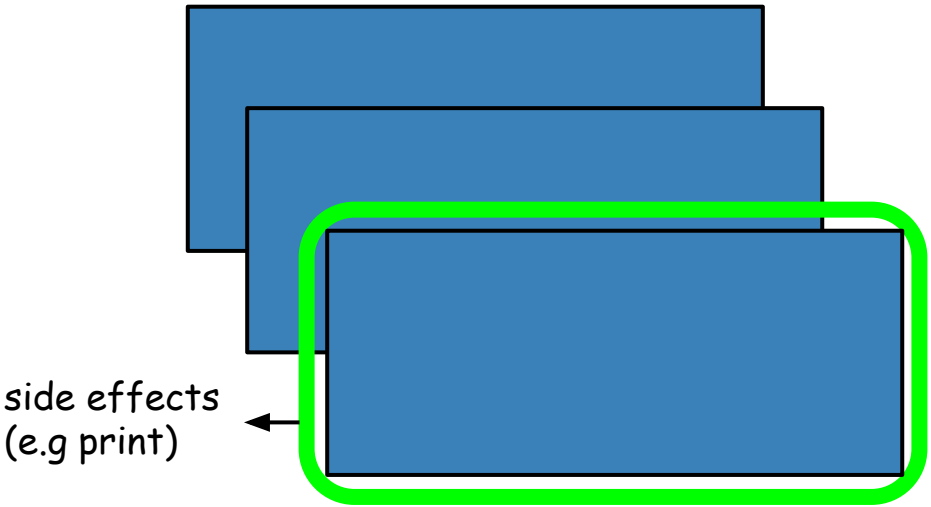

*no return value,
original RDD unchanged

foreachPartition

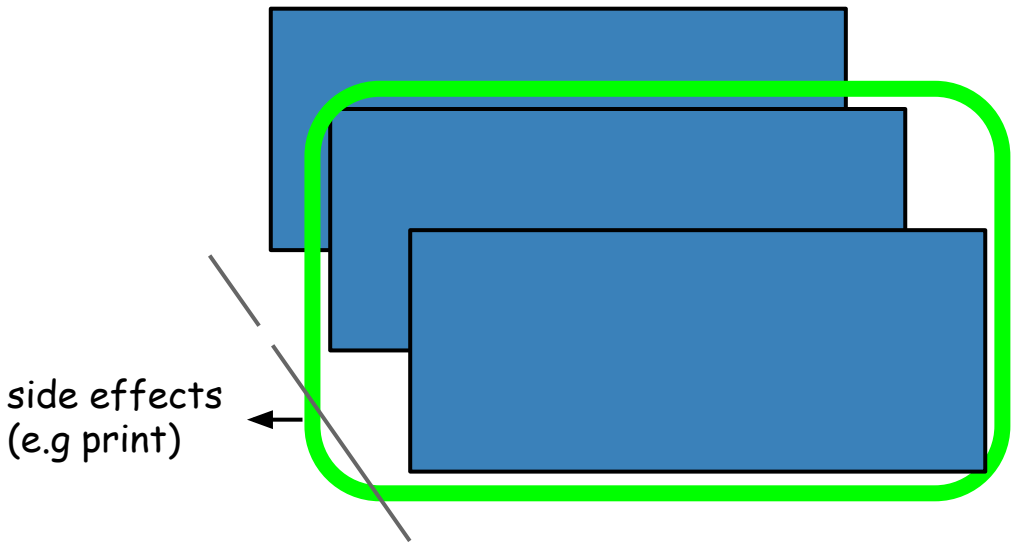

*no return value,
original RDD unchanged

reduce

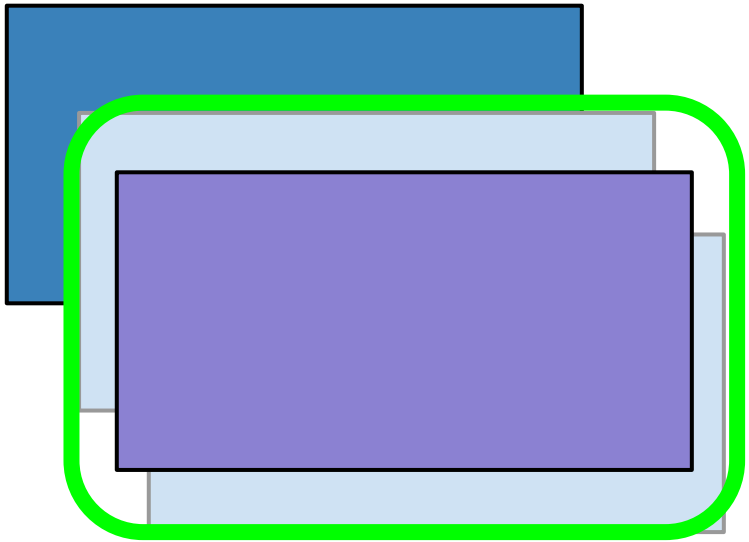

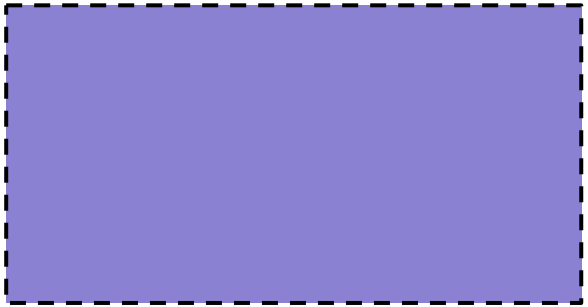

fold

- If you have an empty RDD, it's a replacement to reduce.

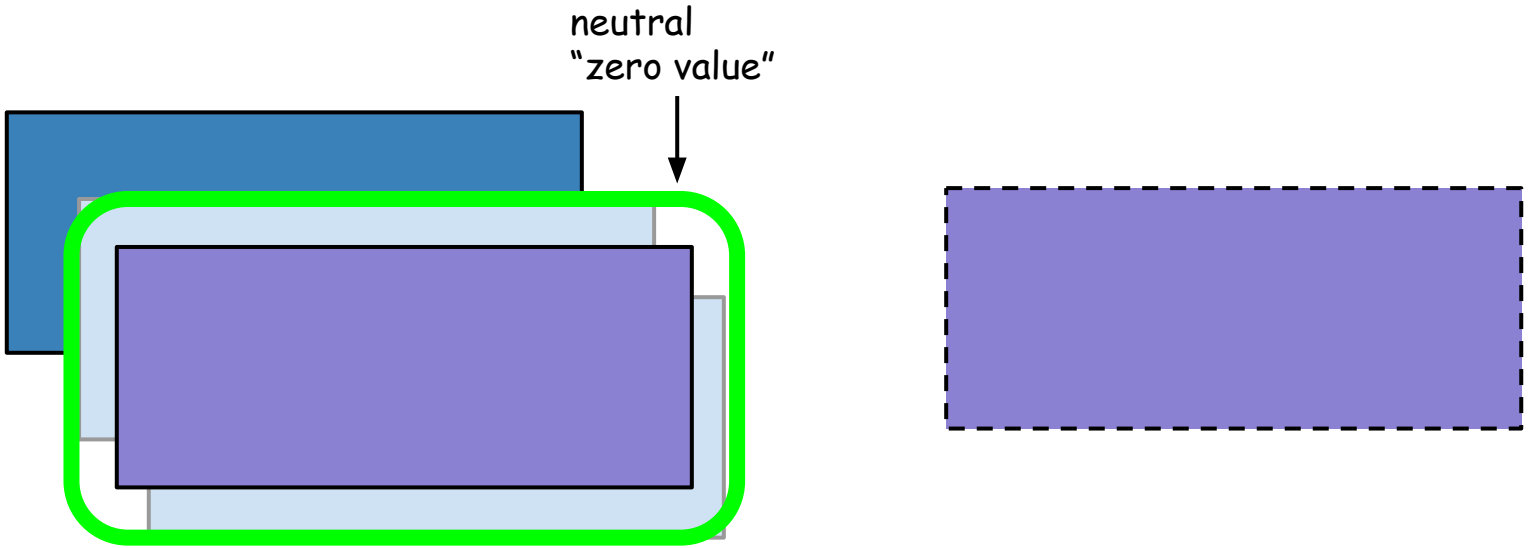

aggregate()

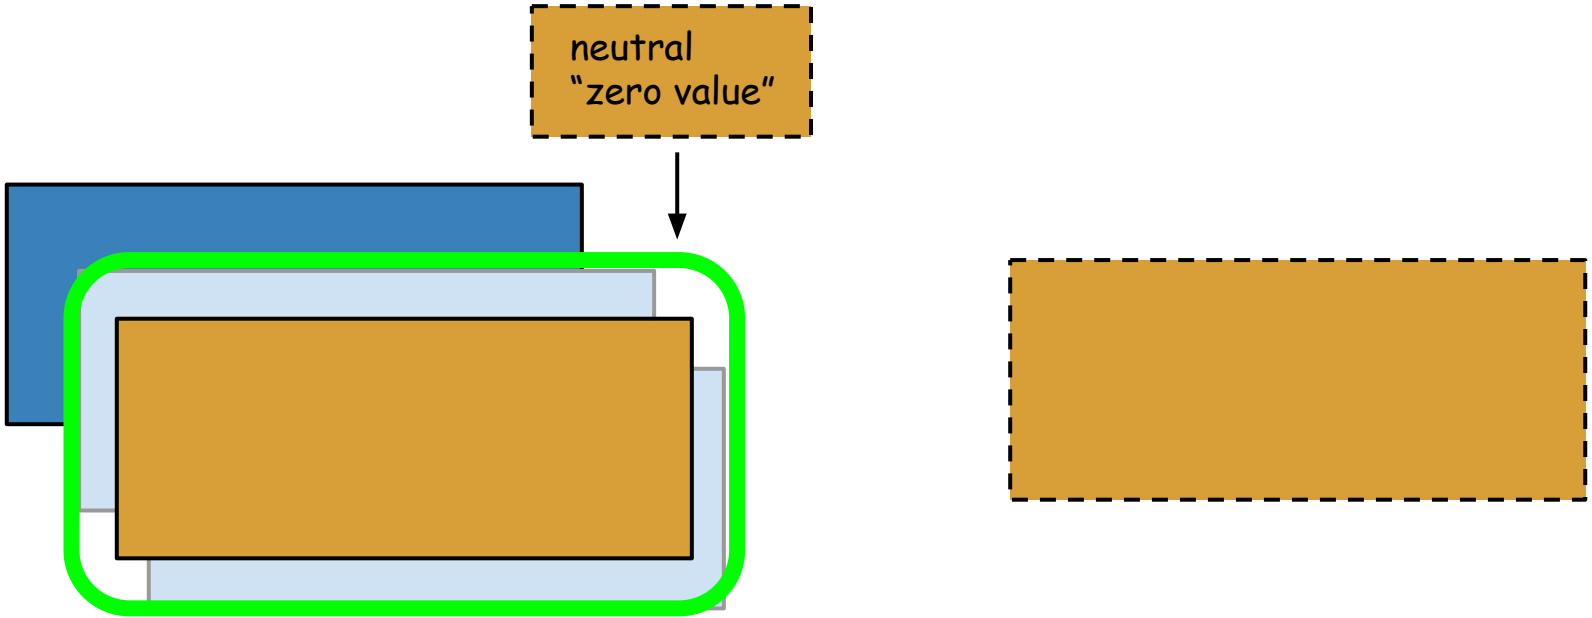

max

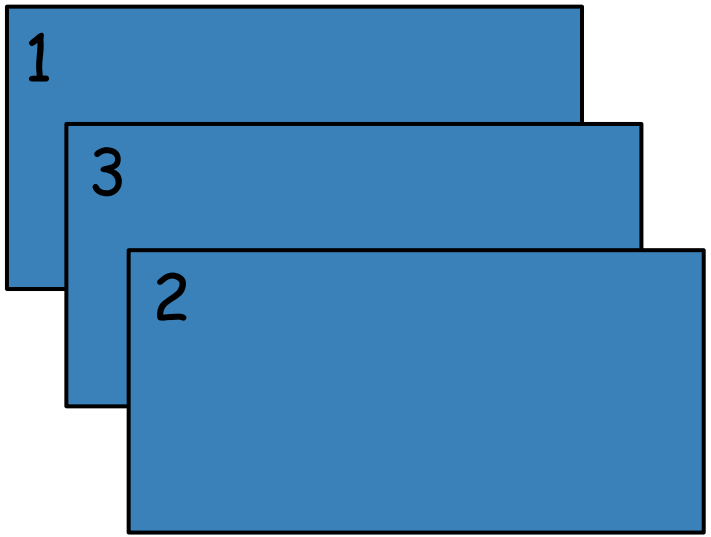

min

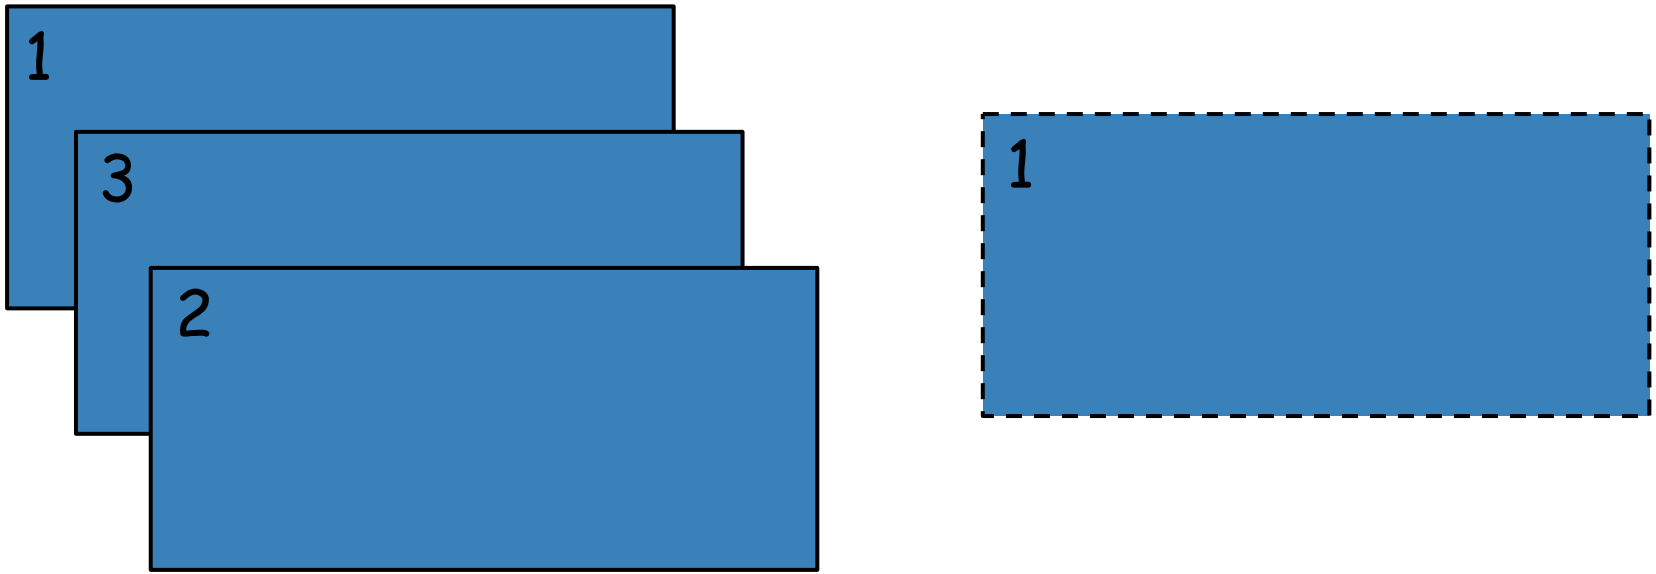

sum

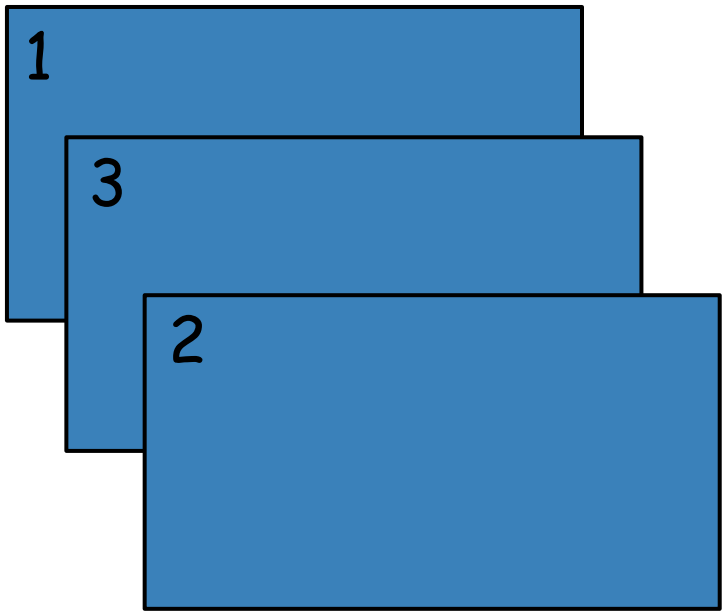

countByValue()

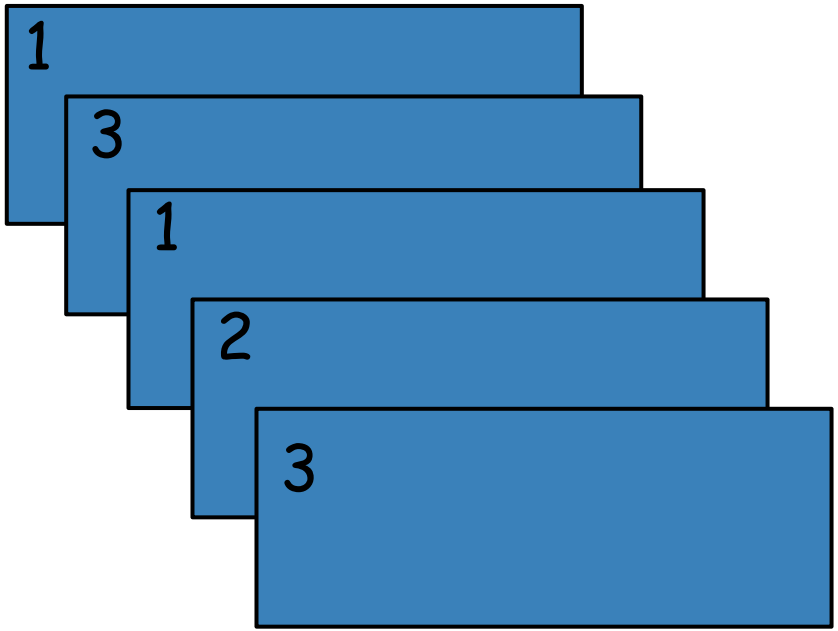

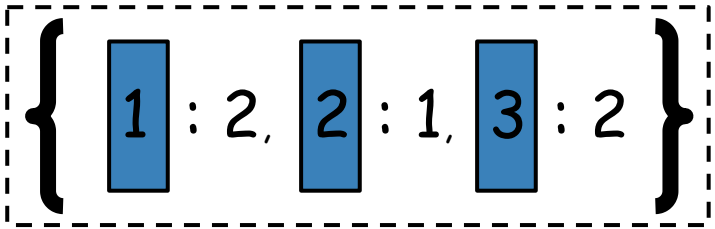

top

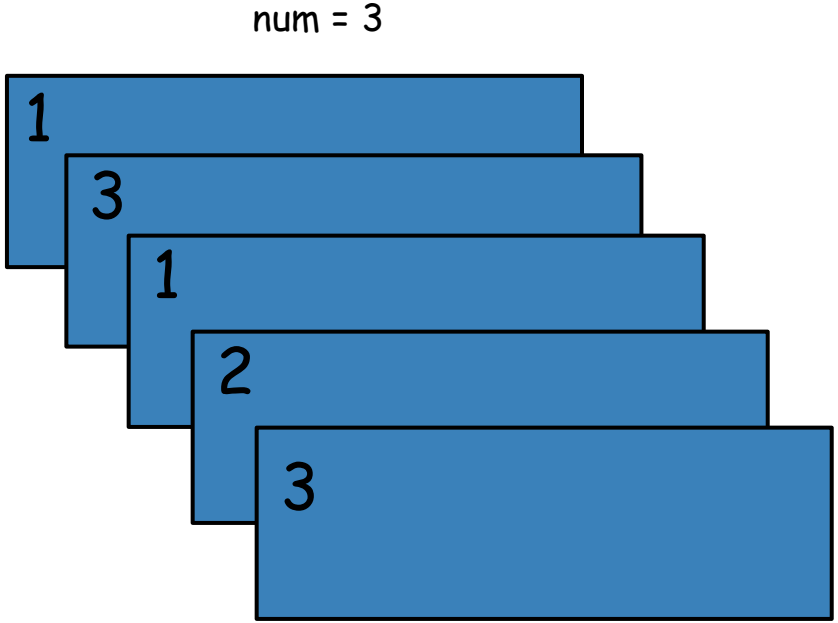

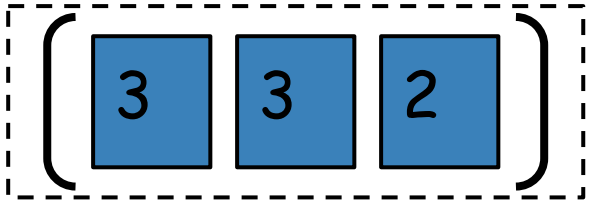

topOrdered

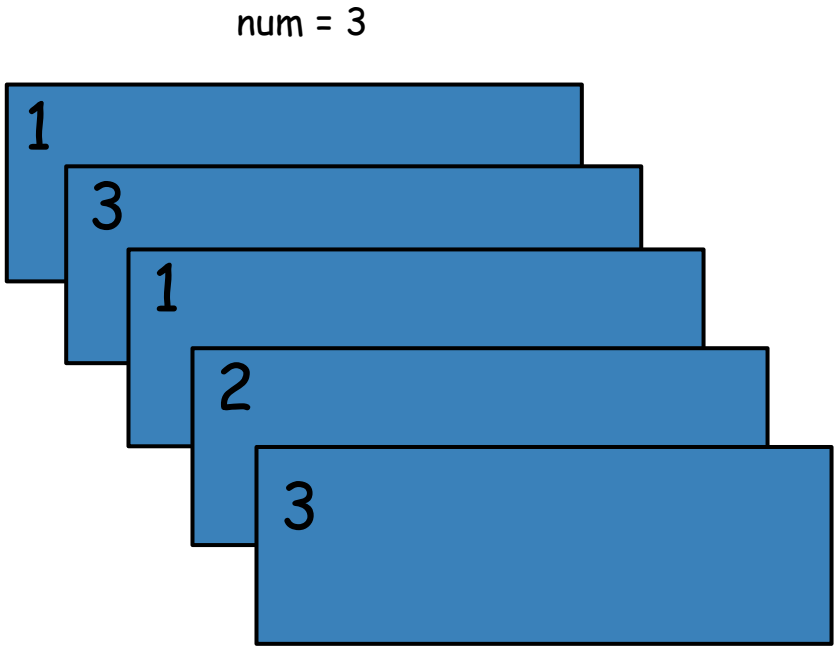

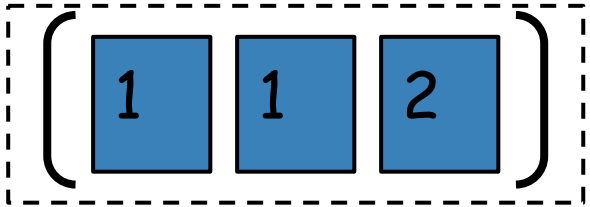

Panorays

take

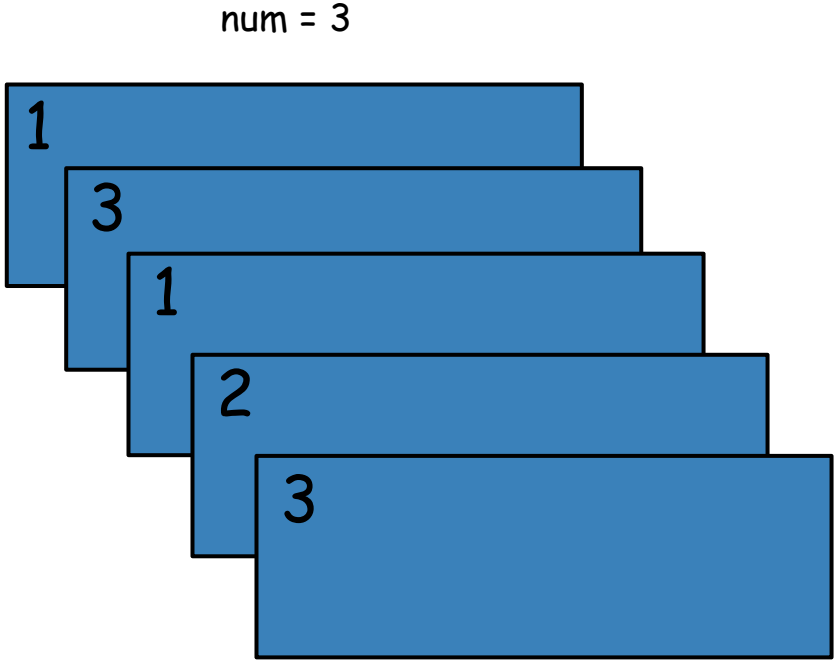

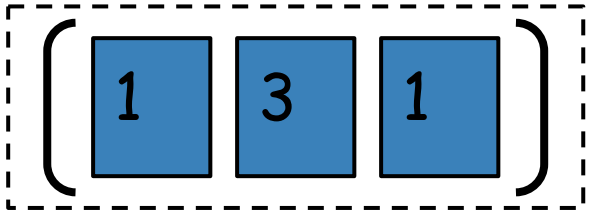

first

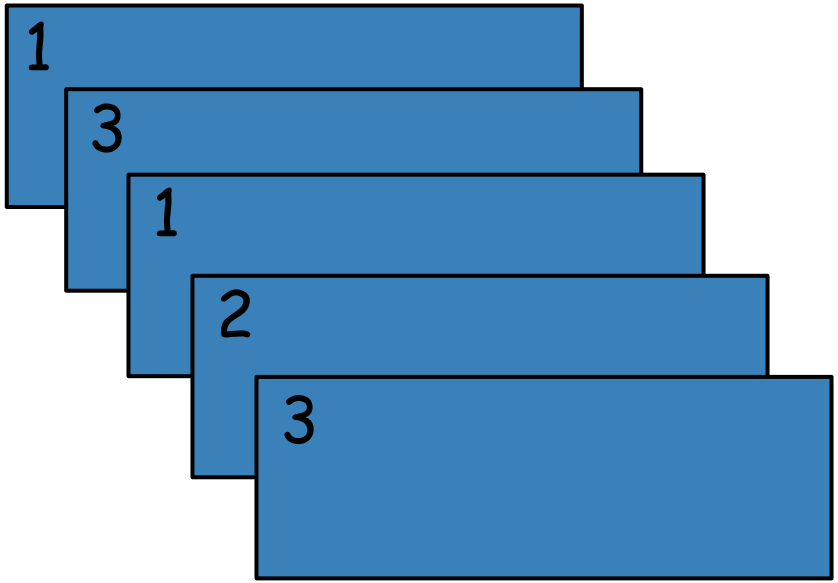

More Actions

- takeSample()
- mean()
- variance()
- tsdev()
- histogram()
- sampleStdev()
- sampleVariance()

Again, to tell you the truth, I've never used these.

Persistence

Actions

- saveAsObjectFile(path)
- saveAsTextFile(path)
- ExternalConnector
- foreach()
  - forEachPartition()

Key Value Methods

collectAsMap

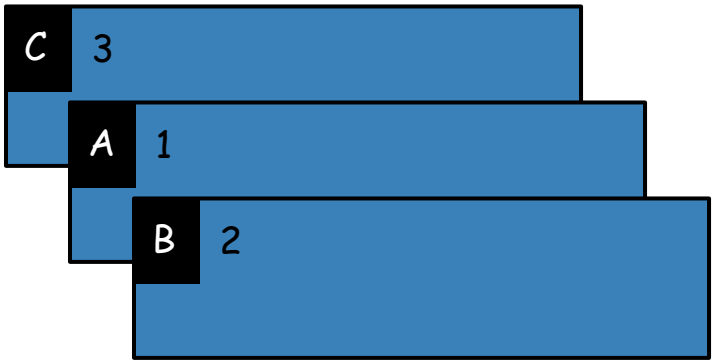

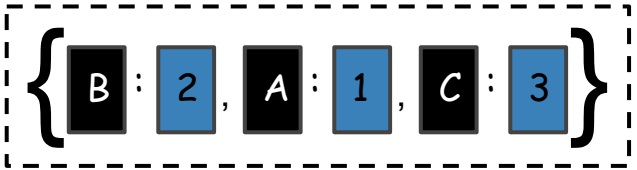

keys

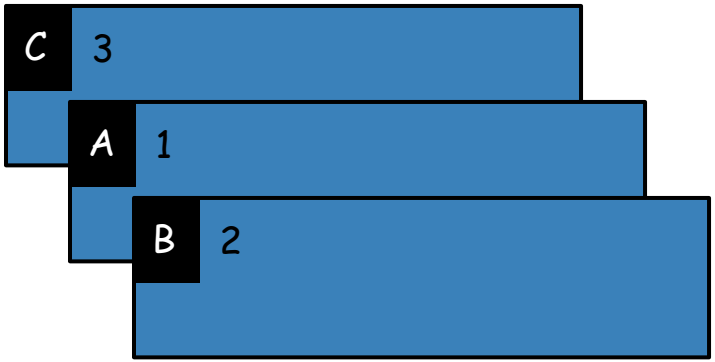

values

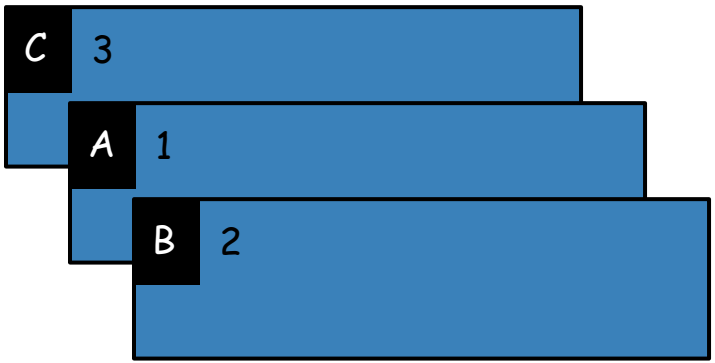

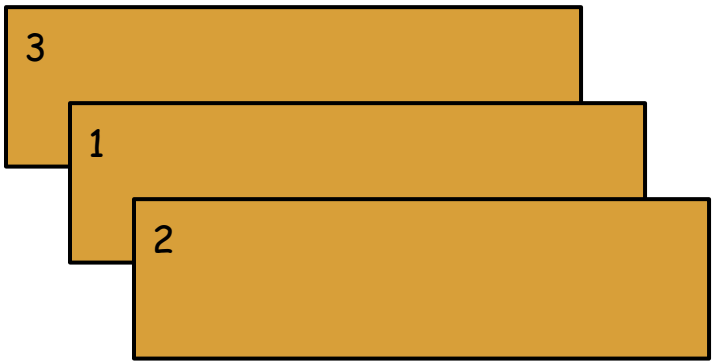

reduceByKey

- Occurs locally

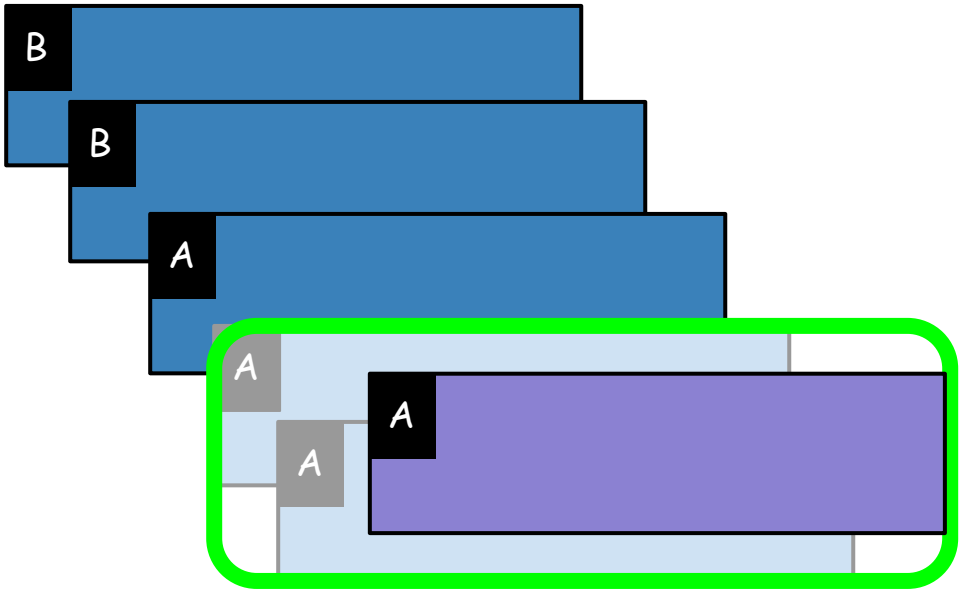

reduceByKeyLocally

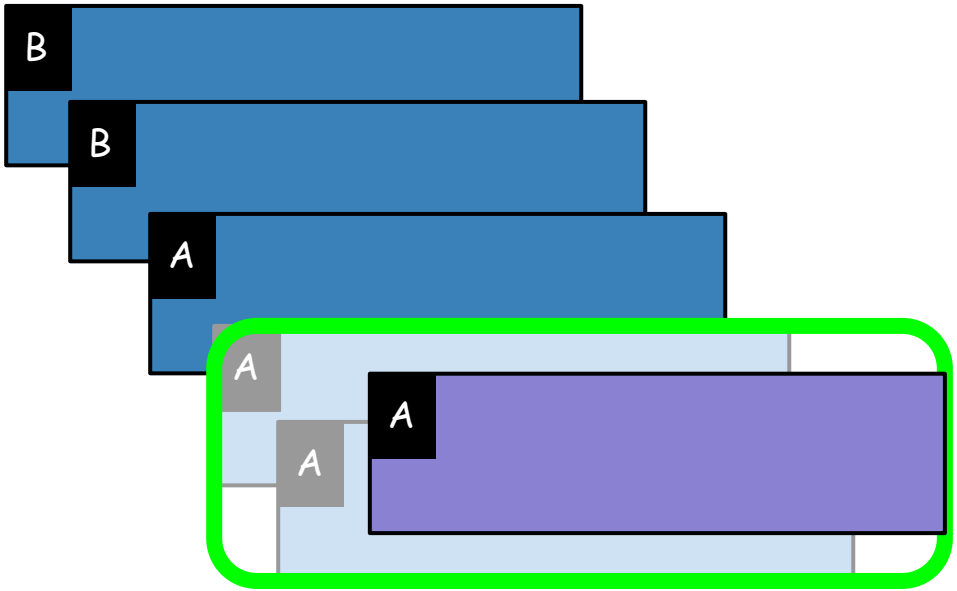

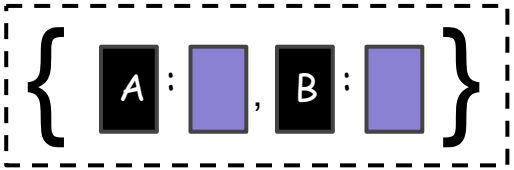

Panorays

countByKey

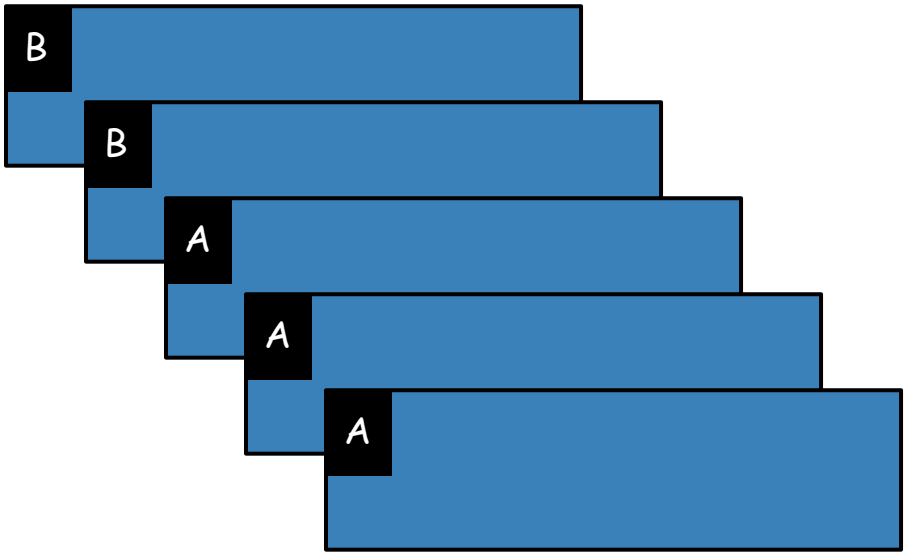

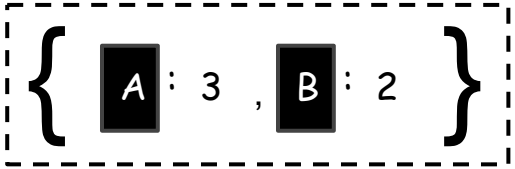

partitionBy

new partition index = key % numPartitions

numPartitions = 3

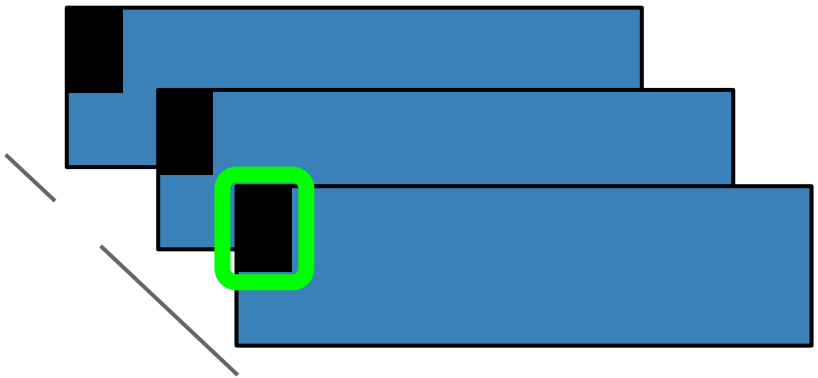

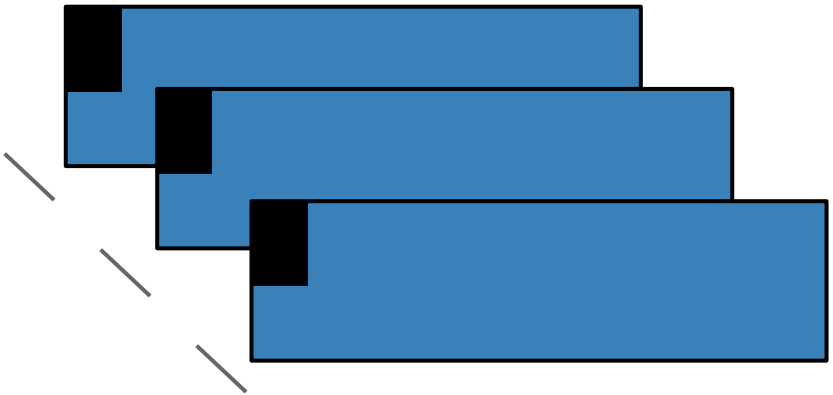

combineBy

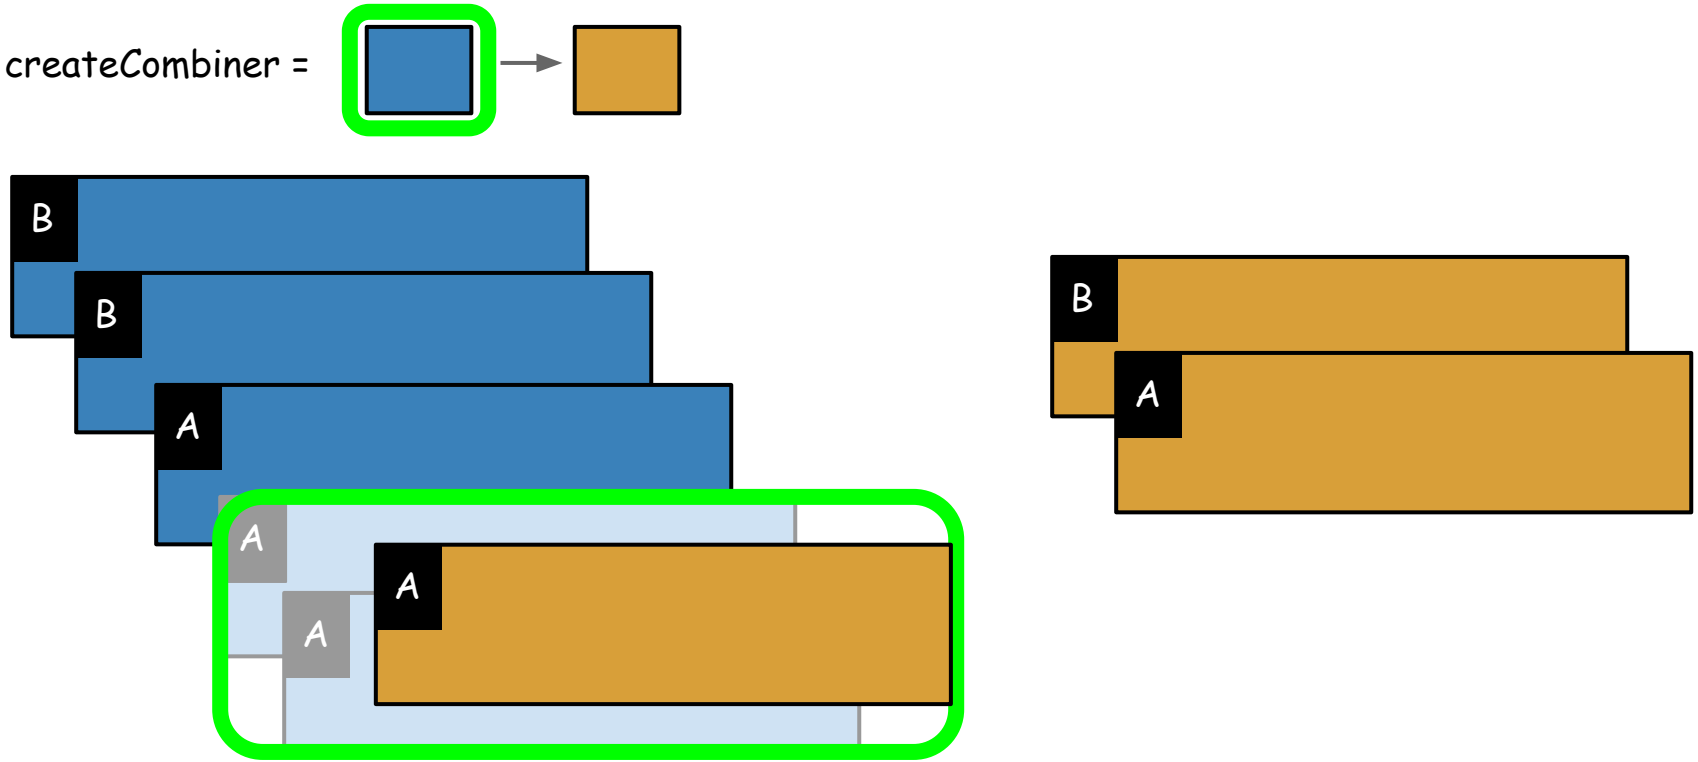

Panorays

aggregateByKey

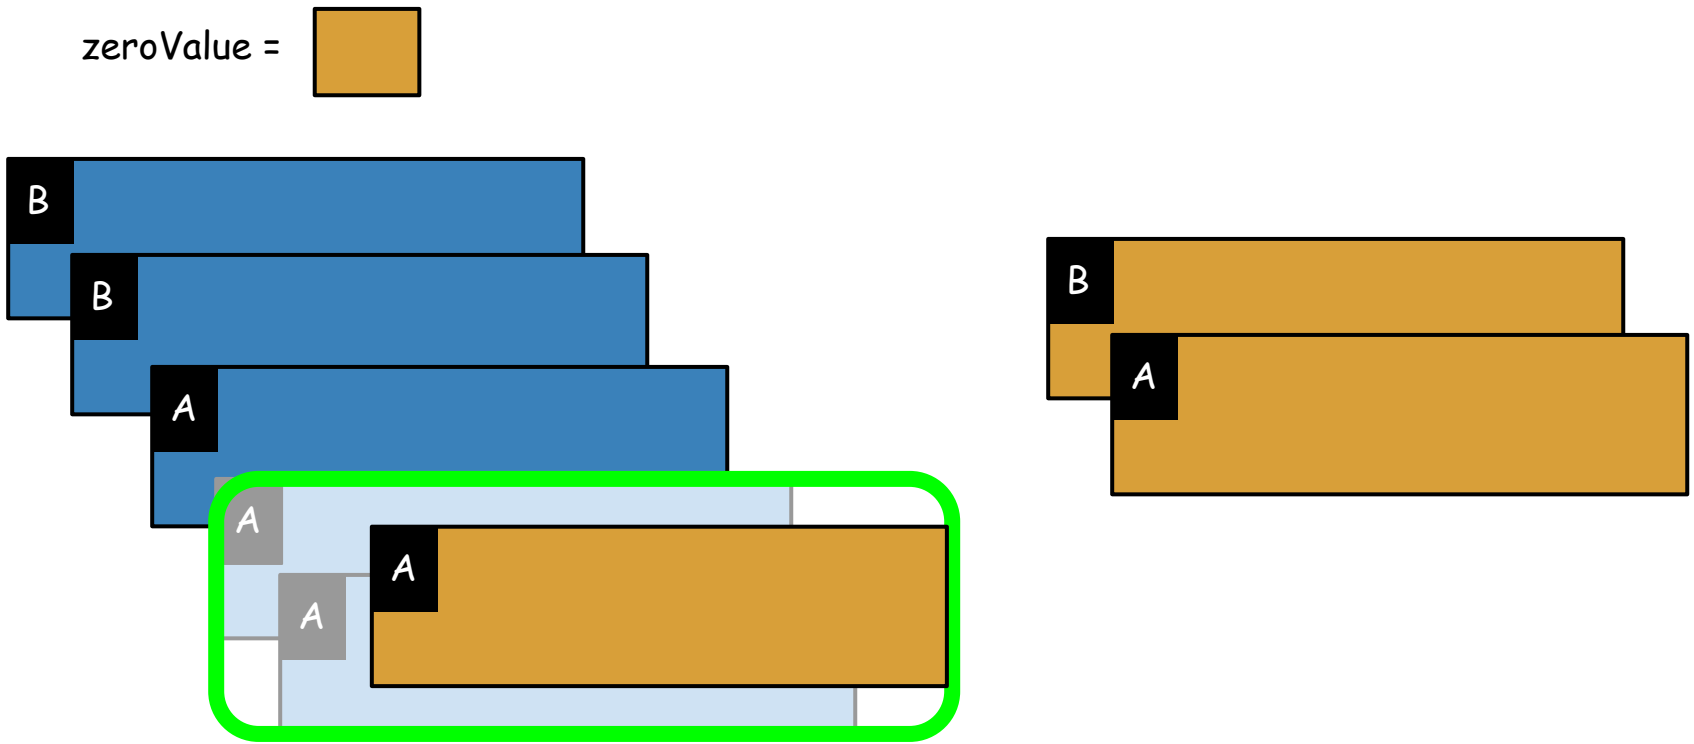

Panorays

foldByKey

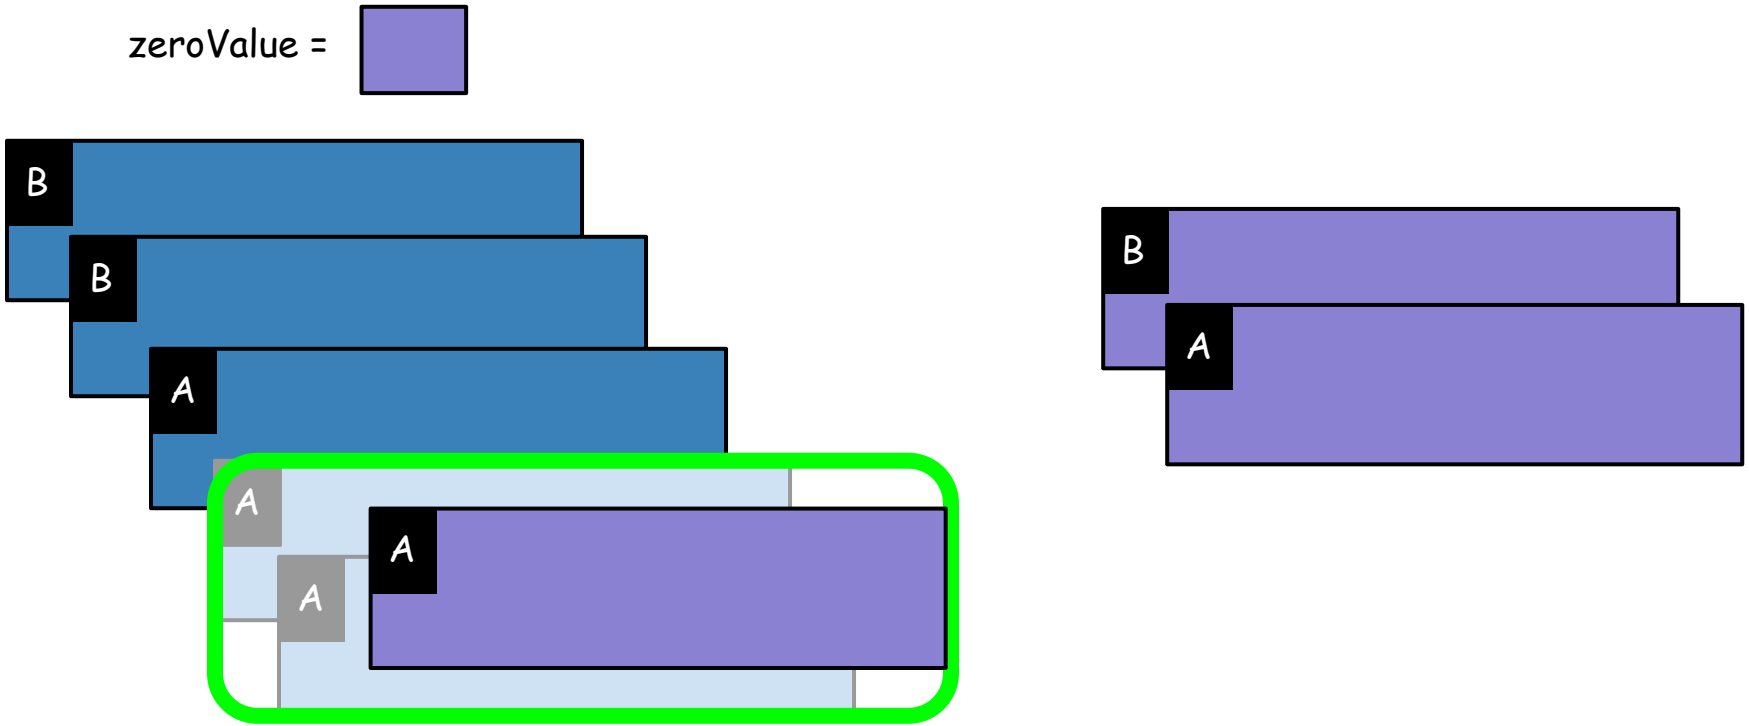

Panorays

Regular Methods

- collectAsMap
- mapValues
- reduceByKey
  - Occurs locally
- foldByKey
- aggregateByKey
- combineByKey
- groupByKey
- countByKey
- sampleByKey
- substractByKey
- sortByKey

SQL-like Pairings

join

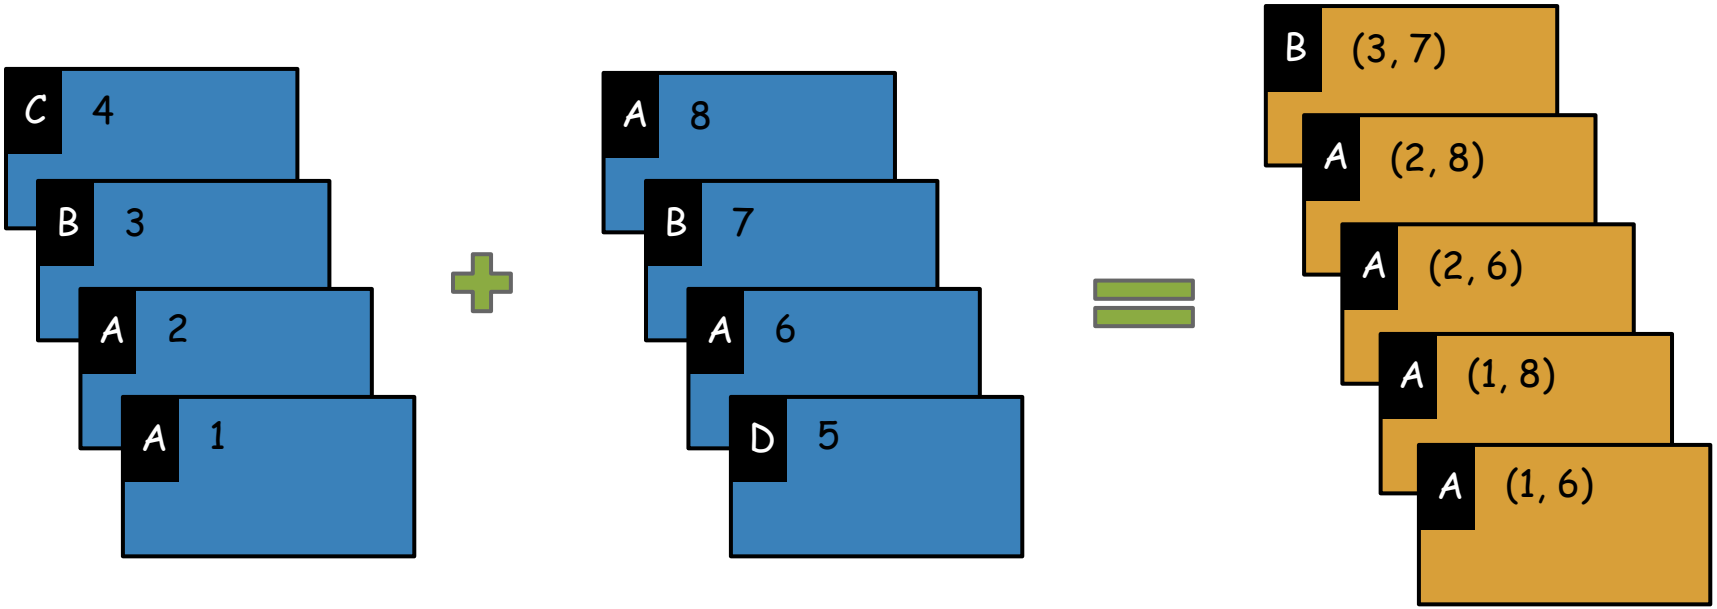

Panorays

fullOuterJoin

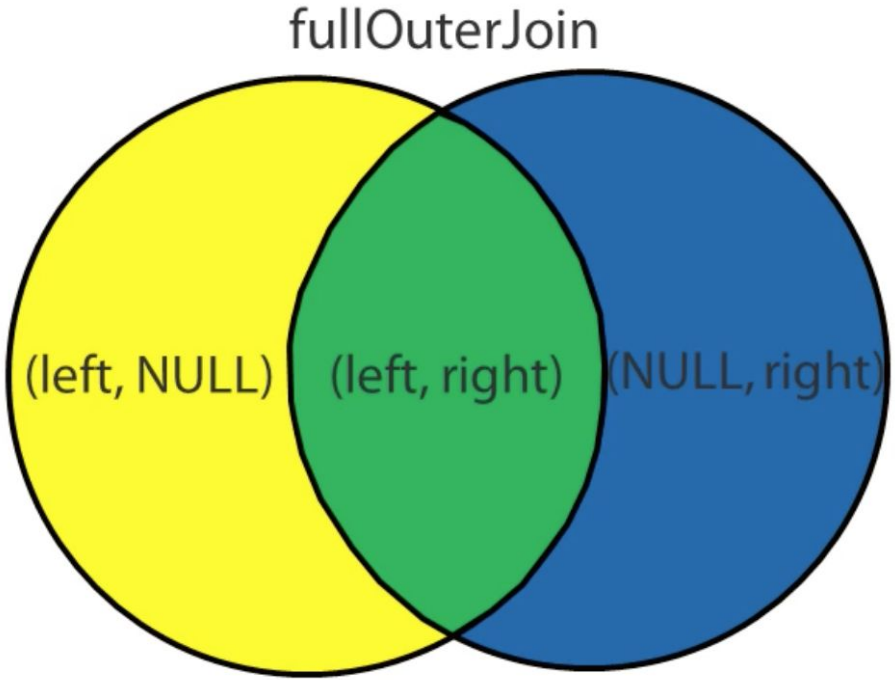

Panorays

leftJoin

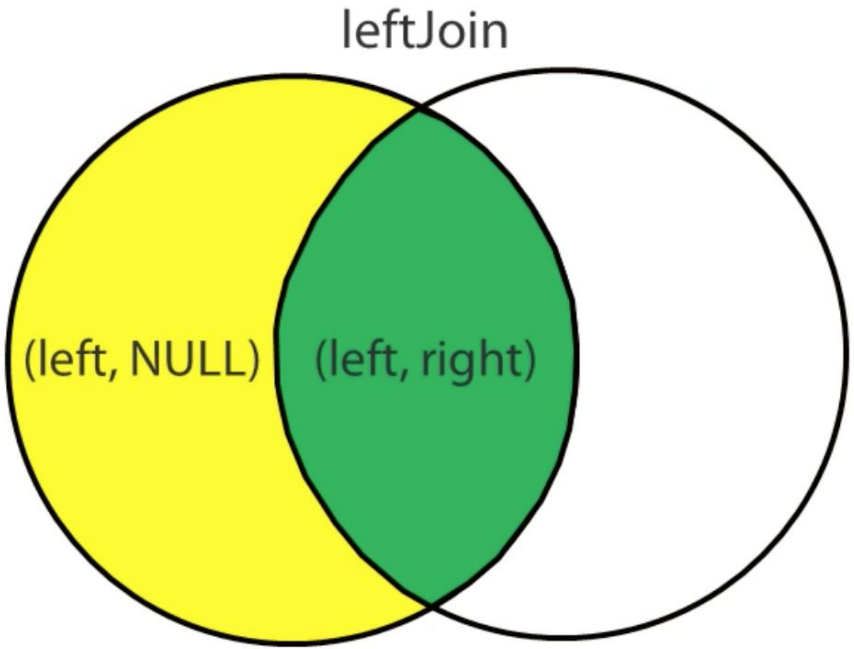

Panorays

rightJoin

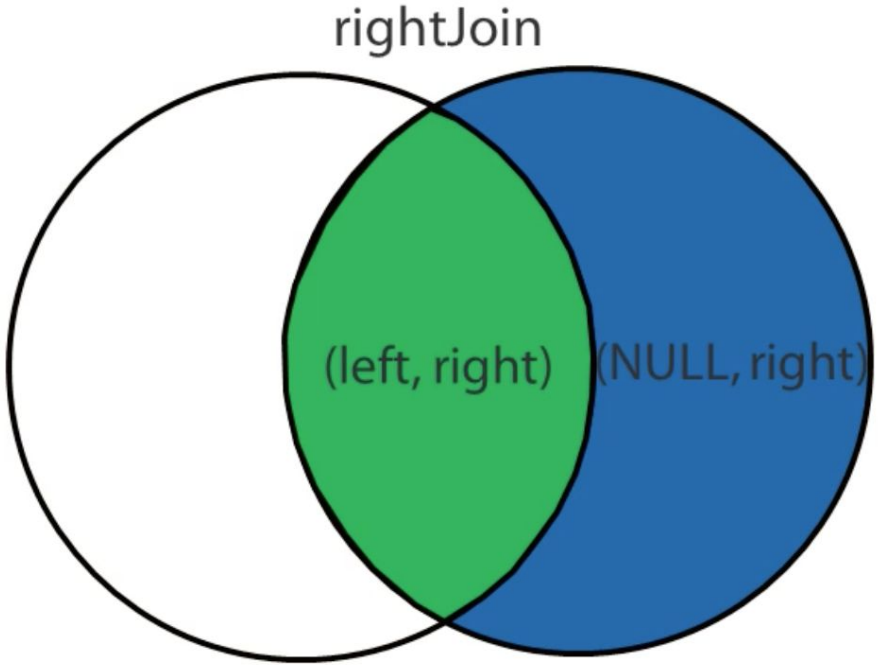

Panorays

cogroup

- When called on datasets of type (K, V) and (K, W), returns a dataset of (K, (Iterable<V>, Iterable<W>)) tuples. This operation is also called groupWith.

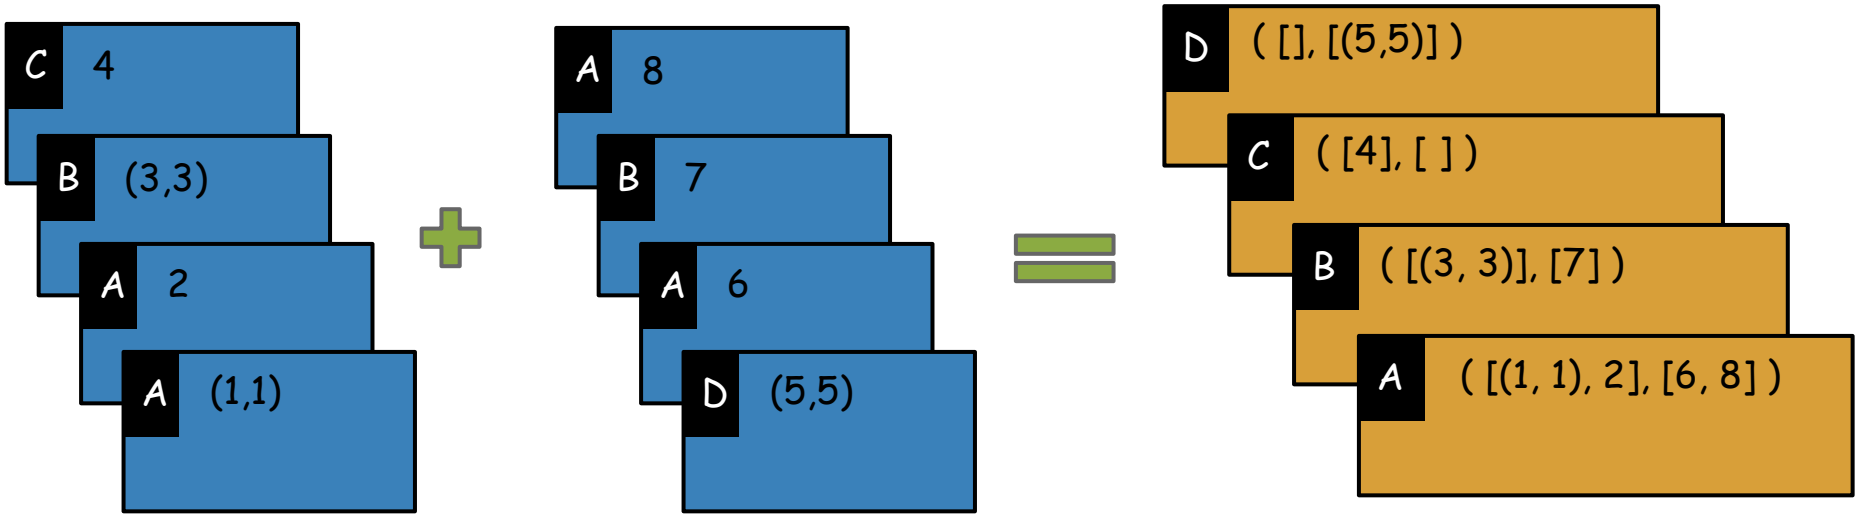

Panorays

List of Transformations and Actions

General

- map
- filter
- flatMap
- mapPartitions
- mapPartitionsWithIndex
- groupBy
- sortBy

Math / Statistical

- sample
- randomSplit

Set Theory / Relational

- union
- intersection
- subtract
- distinct
- cartesian
- zip

Data Structure / I/O

- keyBy
- zipWithIndex
- zipWithUniqueID
- zipPartitions
- coalesce
- repartition
- repartitionAndSortWithinPartitions
- pipe

- reduce
- collect
- aggregate
- fold
- first
- take
- forEach
- top
- treeAggregate
- treeReduce
- forEachPartition
- collectAsMap

- count
- takeSample
- max
- min
- sum
- histogram
- mean
- variance
- stdev
- sampleVariance
- countApprox
- countApproxDistinct

- takeOrdered

- saveAsTextFile
- saveAsSequenceFile
- saveAsObjectFile
- saveAsHadoopDataset
- saveAsHadoopFile
- saveAsNewAPIHadoopDataset
- saveAsNewAPIHadoopFile

Caching Data

Cache

- Cache / Persist
  - org.apache.spark.storage.StorageLevel.MEMORY_ONLY
- persist(newLevel: StorageLevel)
  - MEMORY_ONLY
  - MEMORY_AND_DISK
  - DISK_ONLY
  - MEMORY_ONLY_SER
  - MEMORY_AND_DISK_SER
  - …_2 (Can be replicated to another node)
  - OFF_HEAP
- unpersist(blocking: boolean = true)

Caching pitfalls

- Always unpersist your cached RDDs, but not too soon.
- Try to cache right after wide transformations
  - To avoid the shuffle in case of failure

Accumulator & Broadcast

Accumulator

- val accumulator = sc.accumulator(0, “Accumulator Name”)
- rdd.foreach(x=> {
  doSomething();
  accumulator += 1;
  })

- In the driver:
  - val accumulatorValue = accumulator.value;

Broadcast

```
val indexer = Map(...)

rdd.flatMap(rddVal => indexer.get(rddVal))
```

- What will happen?

Broadcast

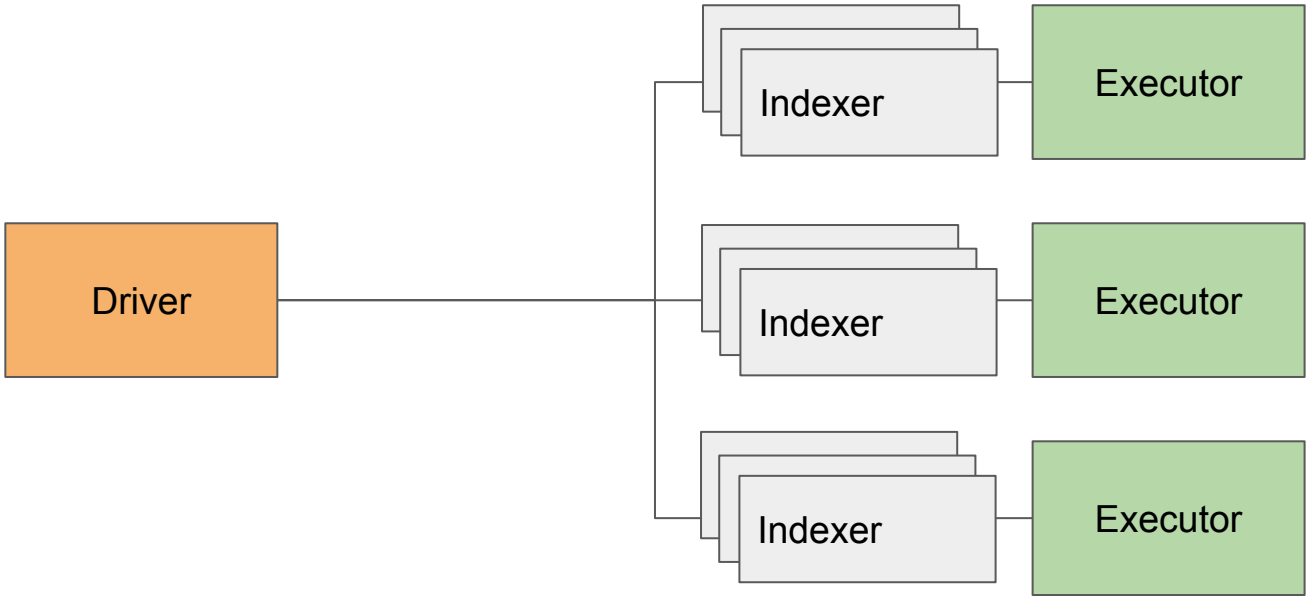
Driver
Indexer
Executor
Indexer
Executor
Indexer
Executor

Broadcast

```
val indexer = sc.broadcast((Map(....))) // Map = 10Mb; indexer < 10Mb

rdd.flatMap(rddVal => indexer.value.get(rddVal))
```

Broadcast

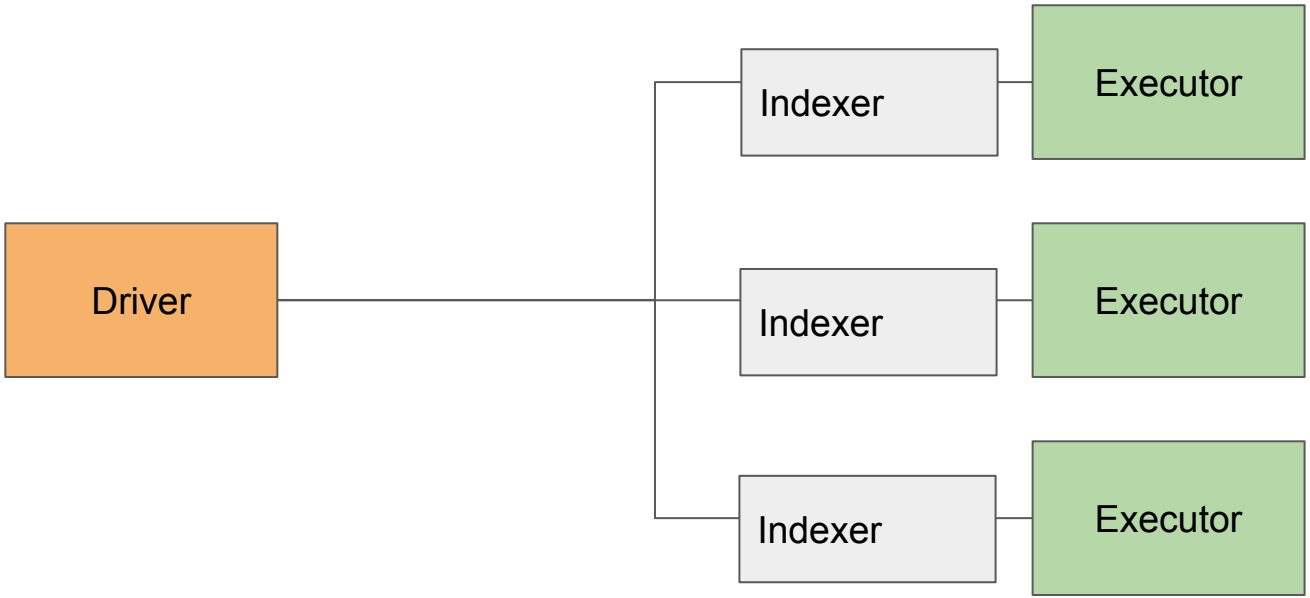
Driver
Indexer
Executor
Indexer
Executor
Indexer
Executor

Java API

- All implemented with:
  - Org.apache.spark.api.java.function
- JavaPairRDD
  - mapToPair
- Almost all of the API is pretty identical, all of the changes can be found in the documentation

Resources

- RDD Research Paper
  - https://www.cs.berkeley.edu/~matei/papers/2012/nsdi_spark.pdf
- Lambdas
  - https://docs.oracle.com/javase/tutorial/java/javaOO/lambdaexpressions.html
- Official Documentation
  - http://spark.apache.org/docs/latest/programming-guide.html

Panorays

Mind the Attack Surface